

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

# Digital Signal Processing

Module 8: Image Processing

- ▶ **Module 8.1:** Introduction to Images and Image Processing
- ▶ **Module 8.2:** Affine Transforms
- ▶ **Module 8.3:** 2D Fourier Analysis
- ▶ **Module 8.4:** Image Filters
- ▶ **Module 8.5:** Image Compression
- ▶ **Module 8.6:** The JPEG Compression Standard

Digital Signal Processing  
Nicola Prandoni and Martin Vetterli  
© 2013

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

# Digital Signal Processing

Module 8.1: Image Processing

- ▶ Images as multidimensional digital signals

- ▶ 2D signal representations

- ▶ Basic signals and operators

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ Images as multidimensional digital signals

- ▶ 2D signal representations

- ▶ Basic signals and operators

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ Images as multidimensional digital signals
- ▶ 2D signal representations
- ▶ Basic signals and operators

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

Please meet ...



- ▶ two-dimensional signal  $x[n_1, n_2]$ ,  $n_1, n_2 \in \mathbb{Z}$

- ▶ indices locate a point on a grid  $\rightarrow$  pixel

- ▶ grid is usually regularly spaced

- ▶ values  $x[n_1, n_2]$  refer to the pixel's appearance

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013



- ▶ two-dimensional signal  $x[n_1, n_2]$ ,  $n_1, n_2 \in \mathbb{Z}$

- ▶ indices locate a point on a grid  $\rightarrow$  pixel

- ▶ grid is usually regularly spaced

- ▶ values  $x[n_1, n_2]$  refer to the pixel's appearance

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ two-dimensional signal  $x[n_1, n_2]$ ,  $n_1, n_2 \in \mathbb{Z}$

- ▶ indices locate a point on a grid  $\rightarrow$  pixel

- ▶ grid is usually regularly spaced

- ▶ values  $x[n_1, n_2]$  refer to the pixel's appearance

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ two-dimensional signal  $x[n_1, n_2]$ ,  $n_1, n_2 \in \mathbb{Z}$
- ▶ indices locate a point on a grid  $\rightarrow$  pixel
- ▶ grid is usually regularly spaced
- ▶ values  $x[n_1, n_2]$  refer to the pixel's appearance

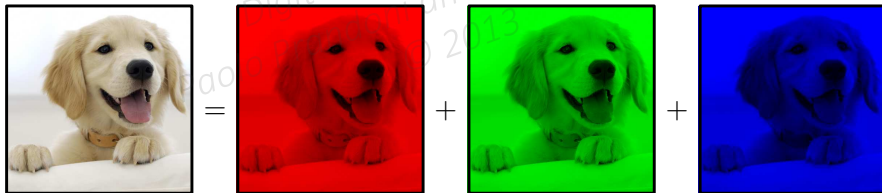
- ▶ grayscale images: scalar pixel values
- ▶ color images: multidimensional pixel values in a color space (RGB, HSV, YUV, etc)
- ▶ we can consider the single components separately:

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ grayscale images: scalar pixel values
- ▶ color images: multidimensional pixel values in a color space (RGB, HSV, YUV, etc)
- ▶ we can consider the single components separately:

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ grayscale images: scalar pixel values
- ▶ color images: multidimensional pixel values in a color space (RGB, HSV, YUV, etc)
- ▶ we can consider the single components separately:



From one to two dimensions...

- ▶ something still works
- ▶ something breaks down
- ▶ something is new

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

From one to two dimensions...

- ▶ something still works
- ▶ something breaks down
- ▶ something is new

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013



From one to two dimensions...

- ▶ something still works
- ▶ something breaks down
- ▶ something is new

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

What works:

- ▶ linearity, convolution

- ▶ Fourier transform

- ▶ interpolation, sampling

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

What works:

- ▶ linearity, convolution

- ▶ Fourier transform

- ▶ interpolation, sampling

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

What works:

- ▶ linearity, convolution
- ▶ Fourier transform
- ▶ interpolation, sampling

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

What breaks down:

- ▶ Fourier analysis less relevant

- ▶ filter design hard, IIRs rare

- ▶ linear operators only mildly useful

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

What breaks down:

- ▶ Fourier analysis less relevant
- ▶ filter design hard, IIRs rare

▶ linear operators only mildly useful

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

What breaks down:

- ▶ Fourier analysis less relevant
- ▶ filter design hard, IIRs rare
- ▶ linear operators only mildly useful

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

What's new:

- ▶ **new manipulations: affine transforms**
- ▶ images are finite-support signals
- ▶ images are (most often) available in their entirety → causality loses meaning
- ▶ images are very specialized signals, designed for a very specific processing system, i.e. the human brain! Lots of semantics that is extremely hard to deal with



What's new:

- ▶ new manipulations: affine transforms
- ▶ images are finite-support signals
- ▶ images are (most often) available in their entirety → causality loses meaning
- ▶ images are very specialized signals, designed for a very specific processing system, i.e. the human brain! Lots of semantics that is extremely hard to deal with

What's new:

- ▶ new manipulations: affine transforms
- ▶ images are finite-support signals
- ▶ images are (most often) available in their entirety → causality loses meaning
- ▶ images are very specialized signals, designed for a very specific processing system, i.e. the human brain! Lots of semantics that is extremely hard to deal with

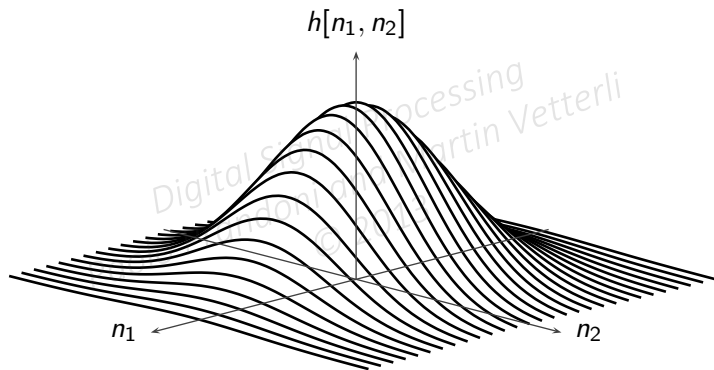
What's new:

- ▶ new manipulations: affine transforms
- ▶ images are finite-support signals
- ▶ images are (most often) available in their entirety → causality loses meaning
- ▶ images are very specialized signals, designed for a very specific processing system, i.e. the human brain! Lots of semantics that is extremely hard to deal with

A two-dimensional discrete-space signal:

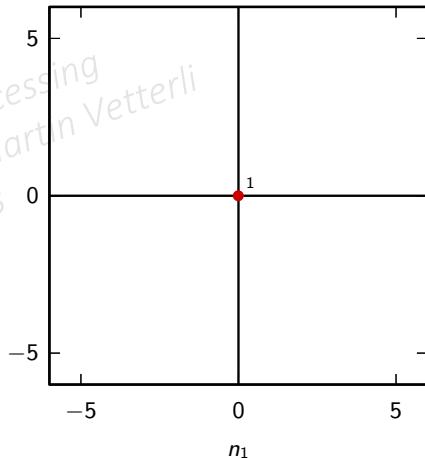
$$x[n_1, n_2], \quad n_1, n_2 \in \mathbb{Z}$$

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

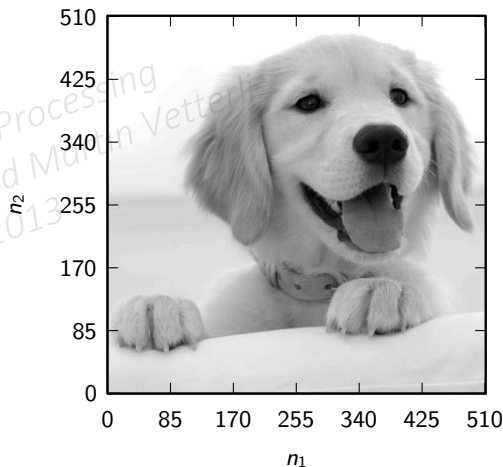


- ▶ just show coordinates of nonzero samples
- ▶ amplitude may be written along explicitly
- ▶ example:

$$\delta[n_1, n_2] = \begin{cases} 1 & \text{if } n_1 = n_2 = 0 \\ 0 & \text{otherwise.} \end{cases}$$



- ▶ medium has a certain dynamic range (paper, screen)
- ▶ image values are quantized (usually to 8 bits, or 256 levels)
- ▶ the eye does the interpolation in space provided the pixel density is high enough



- ▶ images could be unrolled (printers, fax)

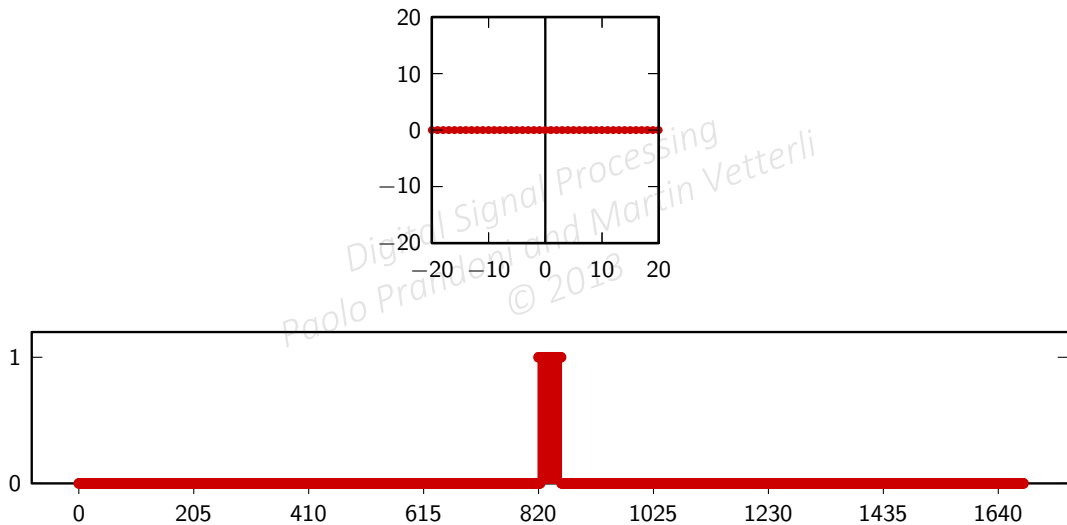
- ▶ but what about spatial correlation?

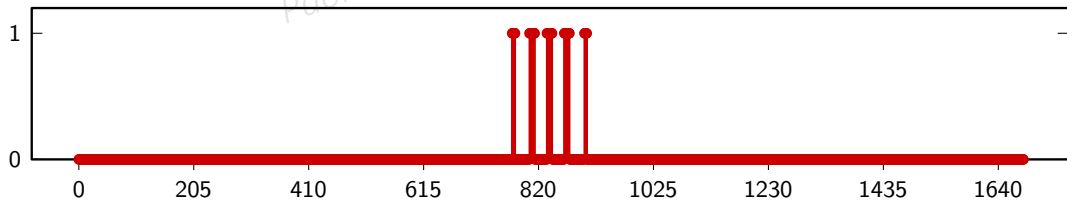
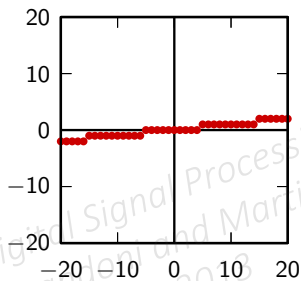
Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

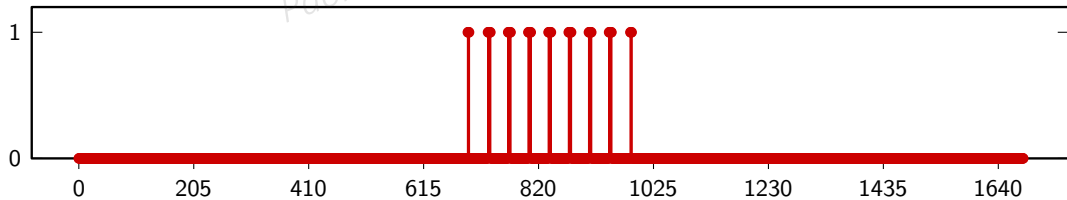
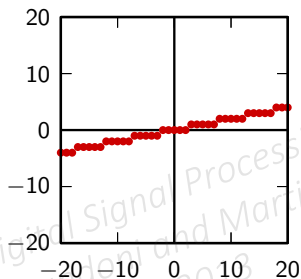


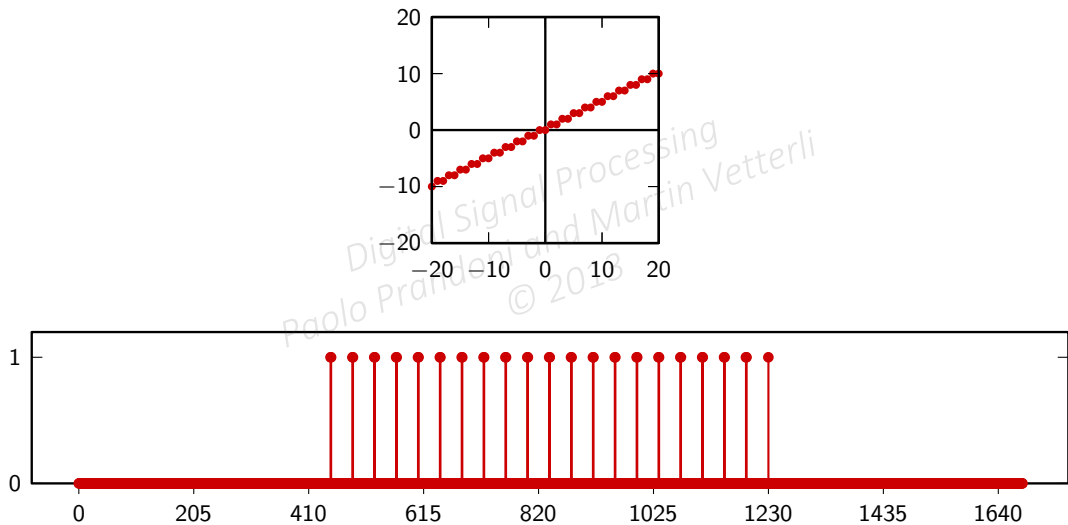
- ▶ images could be unrolled (printers, fax)
- ▶ but what about spatial correlation?

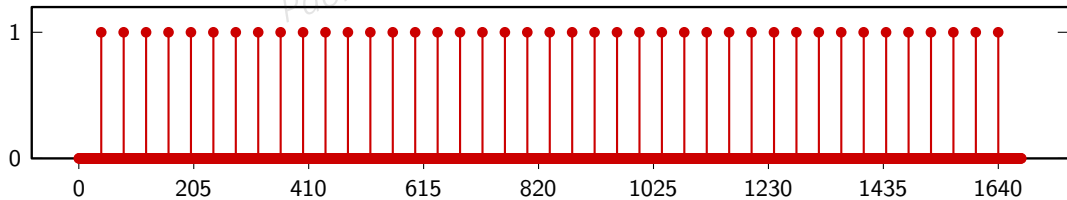
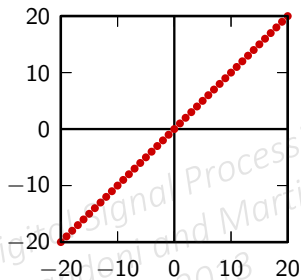
Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

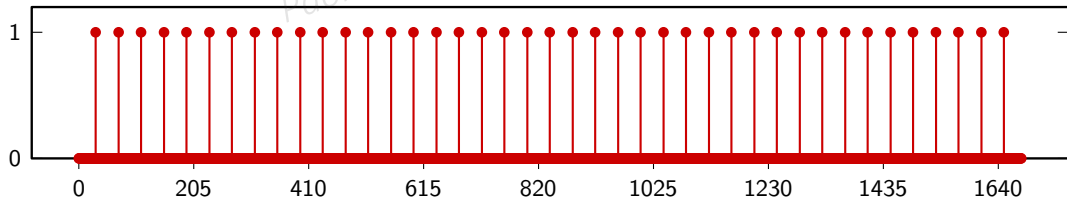
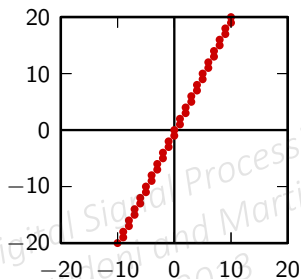


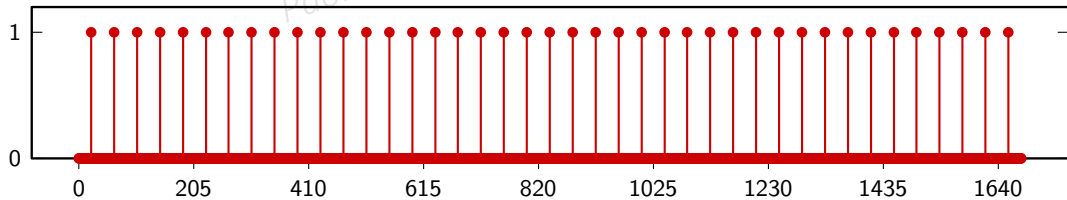
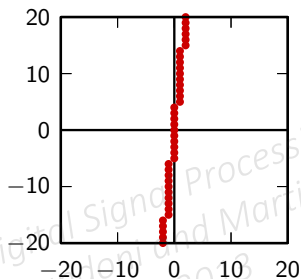




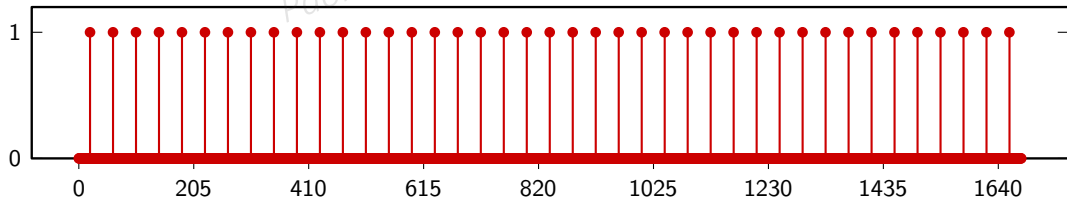
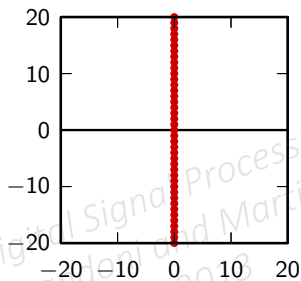




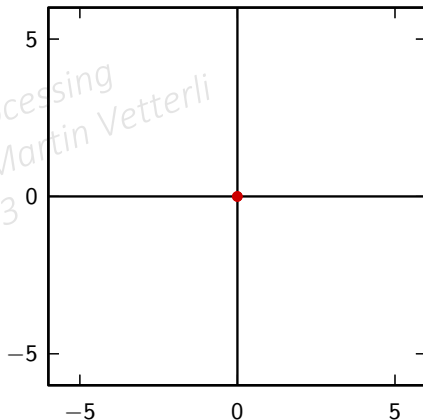




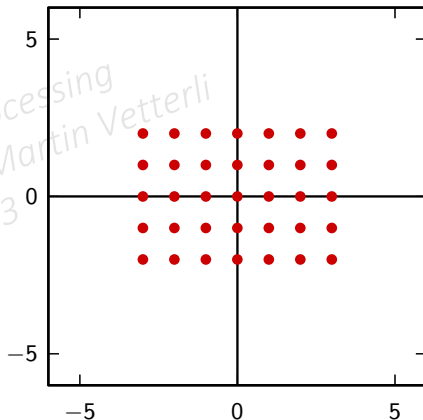




$$\delta[n_1, n_2] = \begin{cases} 1 & \text{if } n_1 = n_2 = 0 \\ 0 & \text{otherwise.} \end{cases}$$



$$\text{rect}\left(\frac{n_1}{2N_1}, \frac{n_2}{2N_2}\right) = \begin{cases} 1 & \text{if } |n_1| < N_1 \\ & \text{and } |n_2| < N_2 \\ 0 & \text{otherwise;} \end{cases}$$



$$x[n_1, n_2] = x_1[n_1]x_2[n_2]$$

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

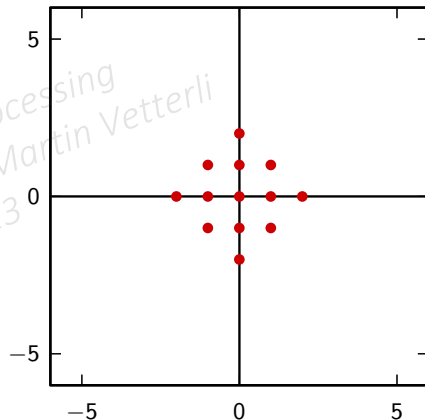
$$\delta[n_1, n_2] = \delta[n_1]\delta[n_2]$$

$$\text{rect}\left(\frac{n_1}{2N_1}, \frac{n_2}{2N_2}\right) = \text{rect}\left(\frac{n_1}{2N_1}\right) \text{rect}\left(\frac{n_2}{2N_2}\right).$$

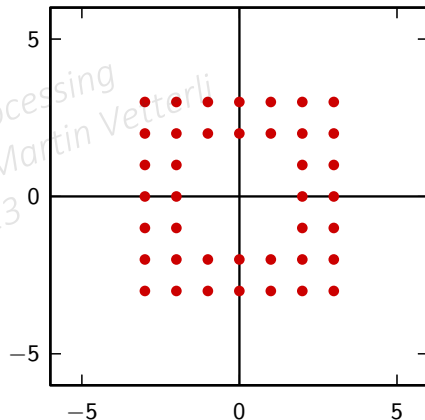
$$\delta[n_1, n_2] = \delta[n_1]\delta[n_2]$$

$$\text{rect}\left(\frac{n_1}{2N_1}, \frac{n_2}{2N_2}\right) = \text{rect}\left(\frac{n_1}{2N_1}\right) \text{rect}\left(\frac{n_2}{2N_2}\right).$$

$$x[n_1, n_2] = \begin{cases} 1 & \text{if } |n_1| + |n_2| \leq N \\ 0 & \text{otherwise} \end{cases}$$



$$x[n_1, n_2] = \text{rect}\left(\frac{n_1}{2N_1}, \frac{n_2}{2N_2}\right) - \text{rect}\left(\frac{n_1}{2M_1}, \frac{n_2}{2M_2}\right)$$





$$x[n_1, n_2] * h[n_1, n_2] = \sum_{k_1=-\infty}^{\infty} \sum_{k_2=-\infty}^{\infty} x[k_1, k_2] h[n_1 - k_1, n_2 - k_2]$$

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

If  $h[n_1, n_2] = h_1[n_1]h_2[n_2]$ :

$$\begin{aligned} x[n_1, n_2] * h[n_1, n_2] &= \sum_{k_1=-\infty}^{\infty} h_1[n_1 - k_1] \sum_{k_2=-\infty}^{\infty} x[k_1, k_2] h_2[n_2 - k_2] \\ &= h_1[n_1] * (h_2[n_2] * x[n_1, n_2]). \end{aligned}$$

If  $h[n_1, n_2]$  is an  $M_1 \times M_2$  finite-support signal:

- ▶ non-separable convolution:  $M_1 M_2$  operations per output sample
- ▶ separable convolution:  $M_1 + M_2$  operations per output sample!

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

END OF MODULE 8.1

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

# Digital Signal Processing

Module 8.2: Image Manipulations

- ▶ Affine transforms
- ▶ Bilinear interpolation

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ Affine transforms
- ▶ Bilinear interpolation

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

mapping  $\mathbb{R}^2 \rightarrow \mathbb{R}^2$  that reshapes the coordinate system:

$$\begin{bmatrix} t'_1 \\ t'_2 \end{bmatrix} = \begin{bmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{bmatrix} \begin{bmatrix} t_1 \\ t_2 \end{bmatrix} - \begin{bmatrix} d_1 \\ d_2 \end{bmatrix}$$

$$\begin{bmatrix} t'_1 \\ t'_2 \end{bmatrix} = \mathbf{A} \begin{bmatrix} t_1 \\ t_2 \end{bmatrix} - \mathbf{d}$$



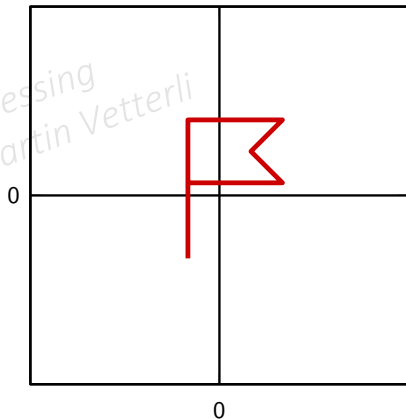
mapping  $\mathbb{R}^2 \rightarrow \mathbb{R}^2$  that reshapes the coordinate system:

$$\begin{bmatrix} t'_1 \\ t'_2 \end{bmatrix} = \begin{bmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{bmatrix} \begin{bmatrix} t_1 \\ t_2 \end{bmatrix} - \begin{bmatrix} d_1 \\ d_2 \end{bmatrix}$$

$$\begin{bmatrix} t'_1 \\ t'_2 \end{bmatrix} = \mathbf{A} \begin{bmatrix} t_1 \\ t_2 \end{bmatrix} - \mathbf{d}$$

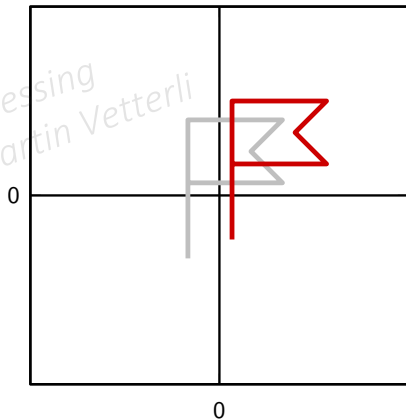
$$\mathbf{A} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} = \mathbf{I}$$

$$\mathbf{d} = \begin{bmatrix} d_1 \\ d_2 \end{bmatrix}$$



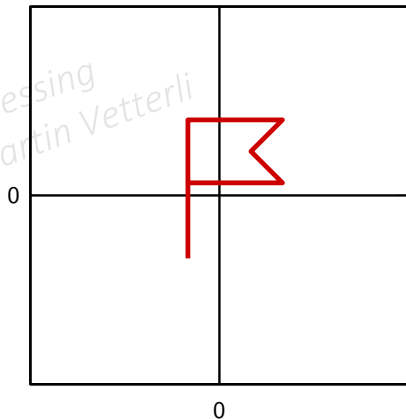
$$\mathbf{A} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} = \mathbf{I}$$

$$\mathbf{d} = \begin{bmatrix} d_1 \\ d_2 \end{bmatrix}$$



$$\mathbf{A} = \begin{bmatrix} a_1 & 0 \\ 0 & a_2 \end{bmatrix}$$
$$\mathbf{d} = 0$$

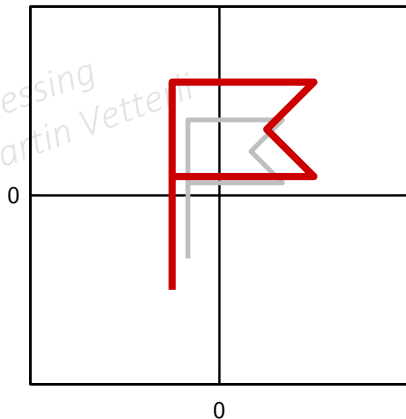
if  $a_1 = a_2$  the *aspect ratio* is preserved



$$\mathbf{A} = \begin{bmatrix} a_1 & 0 \\ 0 & a_2 \end{bmatrix}$$

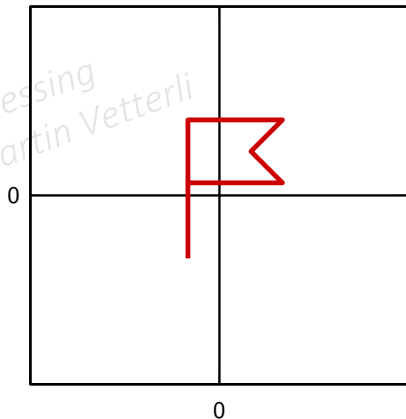
$$\mathbf{d} = 0$$

if  $a_1 = a_2$  the *aspect ratio* is preserved



$$\mathbf{A} = \begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix}$$

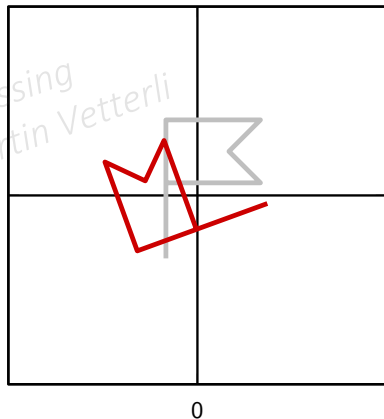
$$\mathbf{d} = 0$$



Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

$$\mathbf{A} = \begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix}$$

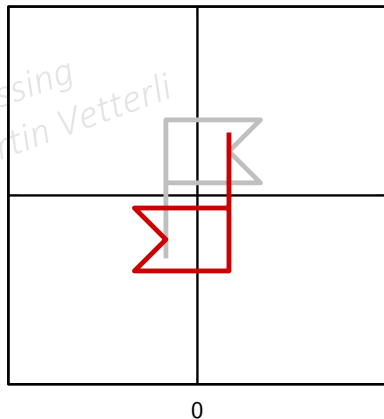
$$\mathbf{d} = \mathbf{0}$$



Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

$$\mathbf{A} = \begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix}$$

$$\mathbf{d} = 0$$





Horizontal:

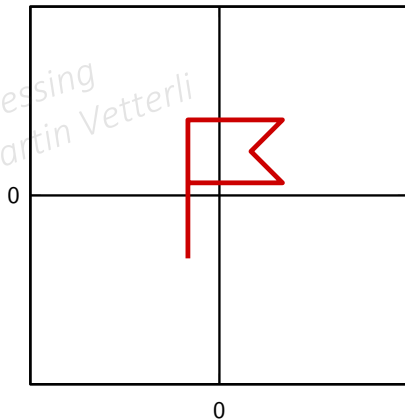
$$\mathbf{A} = \begin{bmatrix} -1 & 0 \\ 0 & 1 \end{bmatrix}$$

$$\mathbf{d} = 0$$

Vertical:

$$\mathbf{A} = \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}$$

$$\mathbf{d} = 0$$



Horizontal:

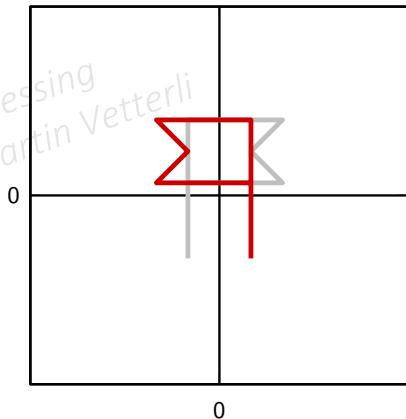
$$\mathbf{A} = \begin{bmatrix} -1 & 0 \\ 0 & 1 \end{bmatrix}$$

$$\mathbf{d} = 0$$

Vertical:

$$\mathbf{A} = \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}$$

$$\mathbf{d} = 0$$



Horizontal:

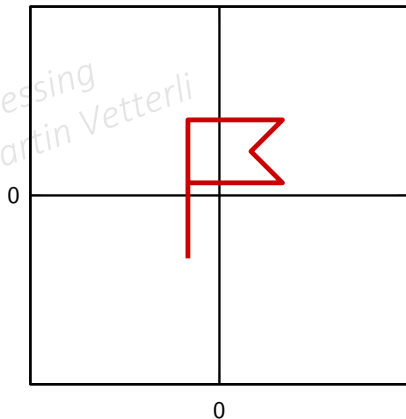
$$\mathbf{A} = \begin{bmatrix} 1 & s \\ 0 & 1 \end{bmatrix}$$

$$\mathbf{d} = 0$$

Vertical:

$$\mathbf{A} = \begin{bmatrix} 1 & 0 \\ s & 1 \end{bmatrix}$$

$$\mathbf{d} = 0$$



Horizontal:

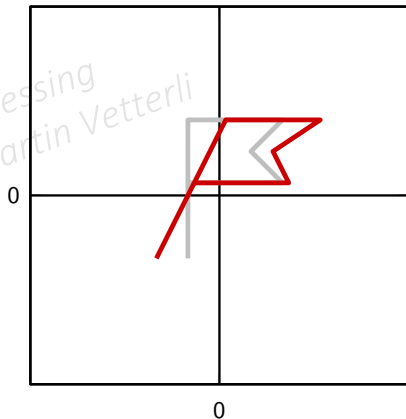
$$\mathbf{A} = \begin{bmatrix} 1 & s \\ 0 & 1 \end{bmatrix}$$

$$\mathbf{d} = 0$$

Vertical:

$$\mathbf{A} = \begin{bmatrix} 1 & 0 \\ s & 1 \end{bmatrix}$$

$$\mathbf{d} = 0$$



$$\begin{bmatrix} t'_1 \\ t'_2 \end{bmatrix} = \mathbf{A} \begin{bmatrix} n_1 \\ n_2 \end{bmatrix} - \mathbf{d} \in \mathbb{R}^2 \neq \mathbb{Z}^2$$

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ apply the *inverse* transform:

$$\begin{bmatrix} t_1 \\ t_2 \end{bmatrix} = \mathbf{A}^{-1} \begin{bmatrix} m_1 + d_1 \\ m_2 + d_2 \end{bmatrix};$$

- ▶ interpolate from the original grid point to the "mid-point"

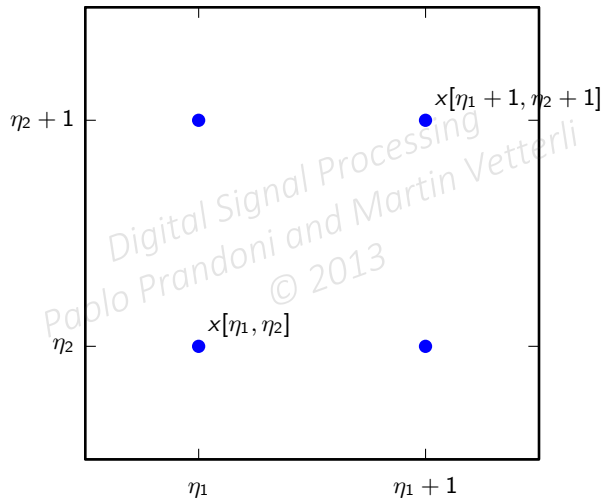
$$(t_1, t_2) = (m_1 + \tau_1, m_2 + \tau_2), \quad \eta_{1,2} \in \mathbb{Z}, \quad 0 \leq \tau_{1,2} < 1$$

- ▶ apply the *inverse* transform:

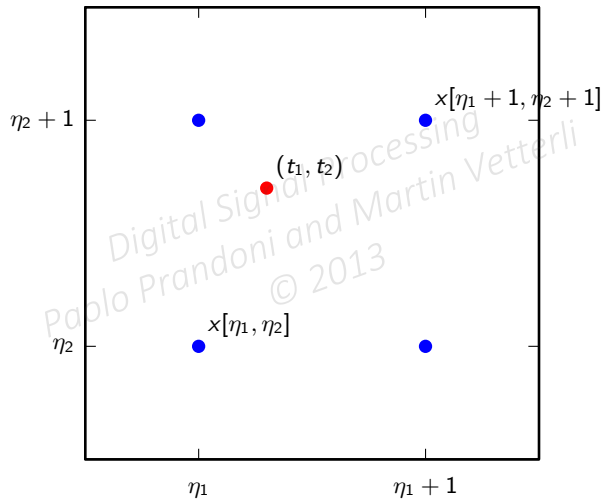
$$\begin{bmatrix} t_1 \\ t_2 \end{bmatrix} = \mathbf{A}^{-1} \begin{bmatrix} m_1 + d_1 \\ m_2 + d_2 \end{bmatrix};$$

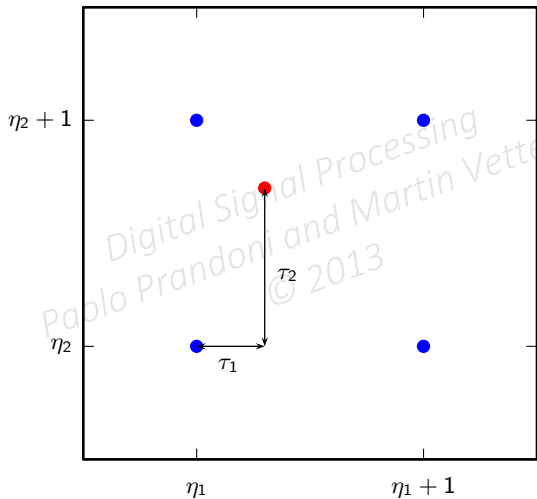
- ▶ interpolate from the original grid point to the “mid-point”

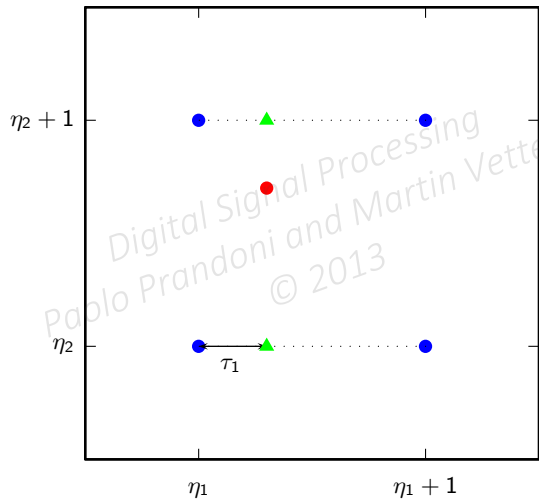
$$(t_1, t_2) = (\eta_1 + \tau_1, \eta_2 + \tau_2), \quad \eta_{1,2} \in \mathbb{Z}, \quad 0 \leq \tau_{1,2} < 1$$

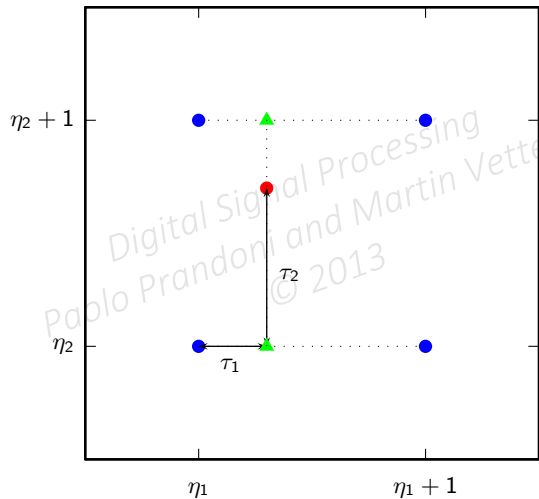












Digital Signal Processing  
Pablo Prandoni and Martin Vetterli  
© 2013

If we use a first-order interpolator:

$$\begin{aligned} y[m_1, m_2] = & (1 - \tau_1)(1 - \tau_2)x[\eta_1, \eta_2] + \tau_1(1 - \tau_2)x[\eta_1 + 1, \eta_2] \\ & + (1 - \tau_1)\tau_2x[\eta_1, \eta_2 + 1] + \tau_1\tau_2x[\eta_1 + 1, \eta_2 + 1] \end{aligned}$$



END OF MODULE 8.2

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

## Digital Signal Processing

Module 8.3: Frequency Analysis



- ▶ DFT

- ▶ Magnitude and phase

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ DFT
- ▶ Magnitude and phase

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

$$X[k_1, k_2] = \sum_{n_1=0}^{N_1-1} \sum_{n_2=0}^{N_2-1} x[n_1, n_2] e^{-j\frac{2\pi}{N_1} n_1 k_1} e^{-j\frac{2\pi}{N_2} n_2 k_2}$$

$$x[n_1, n_2] = \frac{1}{N_1 N_2} \sum_{k_1=0}^{N_1-1} \sum_{k_2=0}^{N_2-1} X[k_1, k_2] e^{j\frac{2\pi}{N_1} n_1 k_1} e^{j\frac{2\pi}{N_2} n_2 k_2}$$

$$X[k_1, k_2] = \sum_{n_1=0}^{N_1-1} \sum_{n_2=0}^{N_2-1} x[n_1, n_2] e^{-j\frac{2\pi}{N_1} n_1 k_1} e^{-j\frac{2\pi}{N_2} n_2 k_2}$$

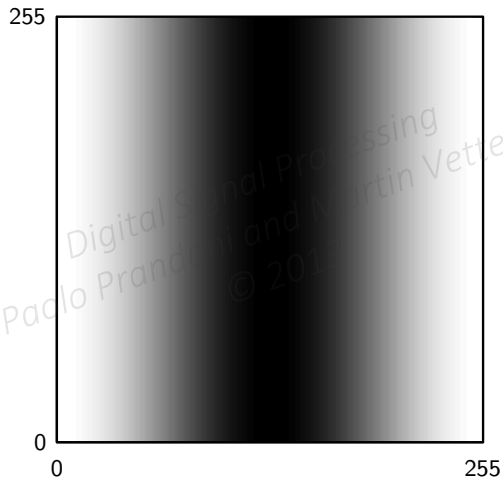
$$x[n_1, n_2] = \frac{1}{N_1 N_2} \sum_{k_1=0}^{N_1-1} \sum_{k_2=0}^{N_2-1} X[k_1, k_2] e^{j\frac{2\pi}{N_1} n_1 k_1} e^{j\frac{2\pi}{N_2} n_2 k_2}$$

There are  $N_1 N_2$  orthogonal basis vectors for an  $N_1 \times N_2$  image:

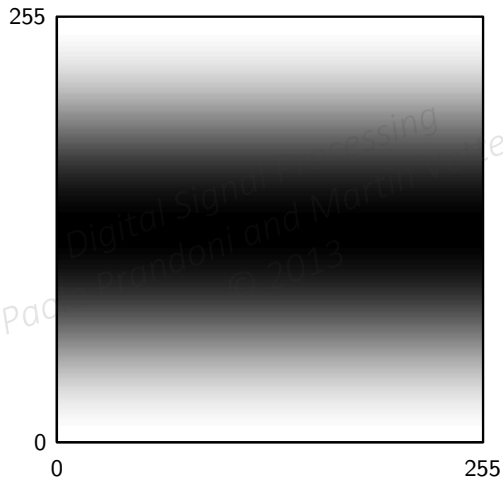
$$w_{k_1, k_2}[n_1, n_2] = e^{j \frac{2\pi}{N_1} n_1 k_1} e^{j \frac{2\pi}{N_2} n_2 k_2}$$

for  $n_1, k_1 = 0, 1, \dots, N_1 - 1$  and  $n_2, k_2 = 0, 1, \dots, N_2 - 1$

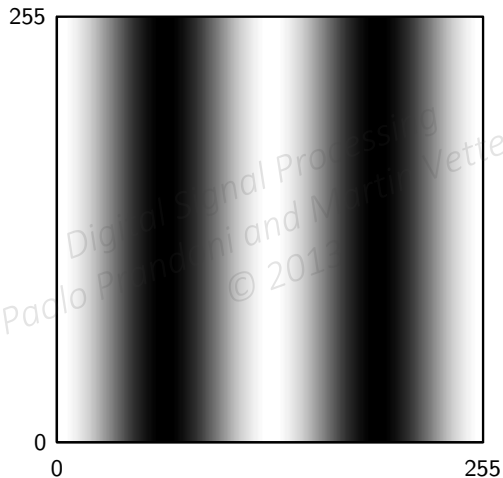
## 2D-DFT basis vectors for $k_1 = 1, k_2 = 0$ (real part)



## 2D-DFT basis vectors for $k_1 = 0, k_2 = 1$ (real part)

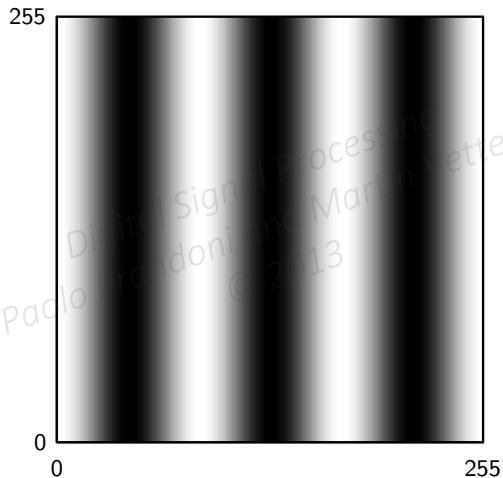


## 2D-DFT basis vectors for $k_1 = 2$ , $k_2 = 0$ (real part)

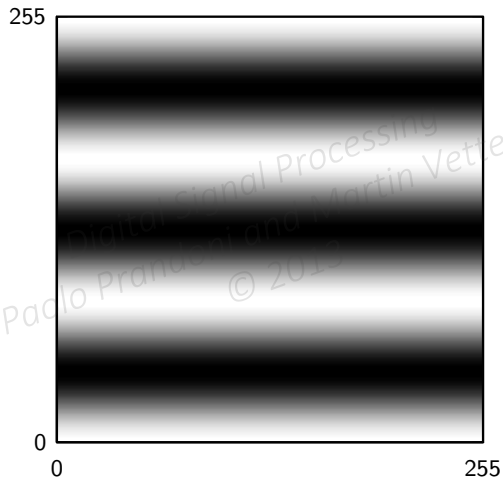




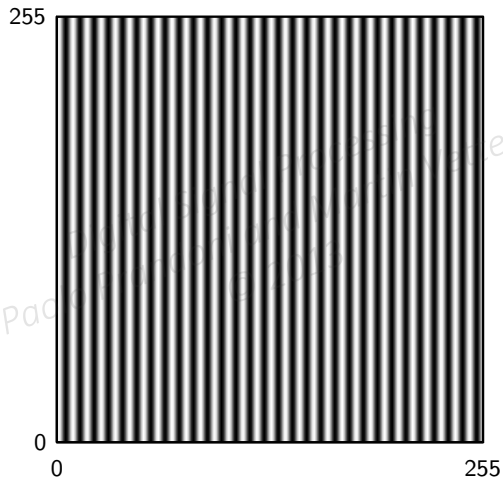
## 2D-DFT basis vectors for $k_1 = 3$ , $k_2 = 0$ (real part)



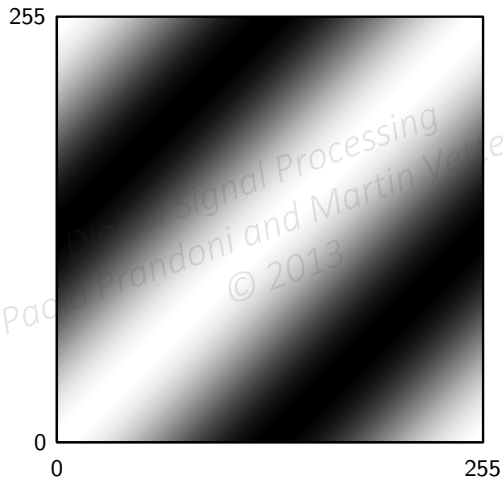
## 2D-DFT basis vectors for $k_1 = 0, k_2 = 3$ (real part)



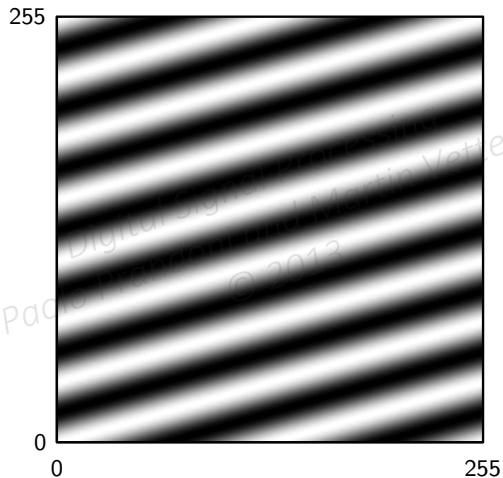
## 2D-DFT basis vectors for $k_1 = 30, k_2 = 0$ (real part)



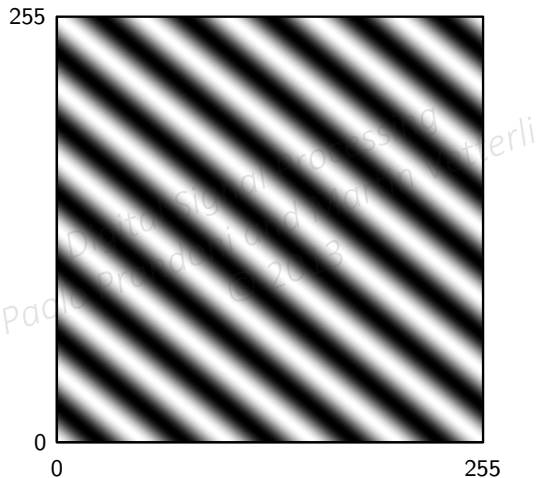
## 2D-DFT basis vectors for $k_1 = 1, k_2 = 1$ (real part)



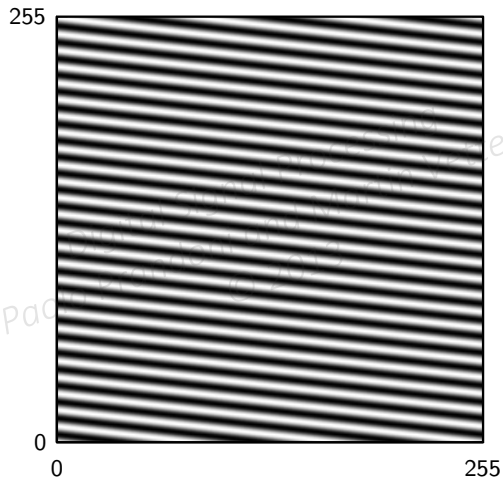
## 2D-DFT basis vectors for $k_1 = 2, k_2 = 7$ (real part)



## 2D-DFT basis vectors for $k_1 = 5$ , $k_2 = 250$ (real part)



## 2D-DFT basis vectors for $k_1 = 3$ , $k_2 = 230$ (real part)



2D-DFT basis functions are separable, and so is the 2D-DFT:

$$X[k_1, k_2] = \sum_{n_1=0}^{N_1-1} \sum_{n_2=0}^{N_2-1} x[n_1, n_2] e^{-j\frac{2\pi}{N_1}n_1k_1} e^{-j\frac{2\pi}{N_2}n_2k_2}$$

► 1D-DFT along  $n_2$  (the columns)

► 1D-DFT along  $n_1$  (the rows)



2D-DFT basis functions are separable, and so is the 2D-DFT:

$$X[k_1, k_2] = \sum_{n_1=0}^{N_1-1} \left[ \sum_{n_2=0}^{N_2-1} x[n_1, n_2] e^{-j\frac{2\pi}{N_2} n_2 k_2} \right] e^{-j\frac{2\pi}{N_1} n_1 k_1}$$

► 1D-DFT along  $n_2$  (the columns)

► 1D-DFT along  $n_1$  (the rows)

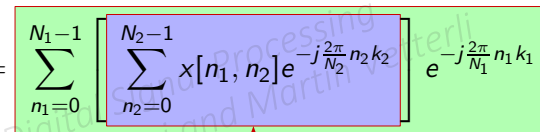
2D-DFT basis functions are separable, and so is the 2D-DFT:

$$X[k_1, k_2] = \sum_{n_1=0}^{N_1-1} \left[ \sum_{n_2=0}^{N_2-1} x[n_1, n_2] e^{-j\frac{2\pi}{N_2} n_2 k_2} \right] e^{-j\frac{2\pi}{N_1} n_1 k_1}$$

► 1D-DFT along  $n_2$  (the columns)

► 1D-DFT along  $n_1$  (the rows)

2D-DFT basis functions are separable, and so is the 2D-DFT:

$$X[k_1, k_2] = \sum_{n_1=0}^{N_1-1} \left[ \sum_{n_2=0}^{N_2-1} x[n_1, n_2] e^{-j\frac{2\pi}{N_2} n_2 k_2} \right] e^{-j\frac{2\pi}{N_1} n_1 k_1}$$


- ▶ 1D-DFT along  $n_2$  (the columns)
- ▶ 1D-DFT along  $n_1$  (the rows)

- ▶ finite-support 2D signal can be written as a matrix  $\mathbf{x}$
- ▶  $N_1 \times N_2$  image is an  $N_2 \times N_1$  matrix ( $n_1$  spans the columns,  $n_2$  spans the rows)
- ▶ recall also the  $N \times N$  DFT matrix (Module 4.2):

$$\mathbf{W}_N = \begin{bmatrix} 1 & 1 & 1 & \dots & 1 \\ 1 & W_N^1 & W_N^2 & \dots & W_N^{N-1} \\ 1 & W_N^2 & W_N^4 & \dots & W_N^{2(N-1)} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & W_N^{N-1} & W_N^{2(N-1)} & \dots & W_N^{(N-1)^2} \end{bmatrix}$$

- ▶ finite-support 2D signal can be written as a matrix  $\mathbf{x}$
- ▶  $N_1 \times N_2$  image is an  $N_2 \times N_1$  matrix ( $n_1$  spans the columns,  $n_2$  spans the rows)
- ▶ recall also the  $N \times N$  DFT matrix (Module 4.2):

$$\mathbf{W}_N = \begin{bmatrix} 1 & 1 & 1 & \dots & 1 \\ 1 & W_N^1 & W_N^2 & \dots & W_N^{N-1} \\ 1 & W_N^2 & W_N^4 & \dots & W_N^{2(N-1)} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & W_N^{N-1} & W_N^{2(N-1)} & \dots & W_N^{(N-1)^2} \end{bmatrix}$$

- ▶ finite-support 2D signal can be written as a matrix  $\mathbf{x}$
- ▶  $N_1 \times N_2$  image is an  $N_2 \times N_1$  matrix ( $n_1$  spans the columns,  $n_2$  spans the rows)
- ▶ recall also the  $N \times N$  DFT matrix (Module 4.2):

$$\mathbf{W}_N = \begin{bmatrix} 1 & 1 & 1 & 1 & \dots & 1 \\ 1 & W_N^1 & W_N^2 & W_N^3 & \dots & W_N^{N-1} \\ 1 & W_N^2 & W_N^4 & W_N^6 & \dots & W_N^{2(N-1)} \\ \vdots & \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & W_N^{N-1} & W_N^{2(N-1)} & W_N^{3(N-1)} & \dots & W_N^{(N-1)^2} \end{bmatrix}$$

$$X[k_1, k_2] = \sum_{n_1=0}^{N_1-1} \left[ \sum_{n_2=0}^{N_2-1} x[n_1, n_2] e^{-j\frac{2\pi}{N_2} n_2 k_2} \right] e^{-j\frac{2\pi}{N_1} n_1 k_1}$$

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

$$X[k_1, k_2] = \sum_{n_1=0}^{N_1-1} \left[ \sum_{n_2=0}^{N_2-1} x[n_1, n_2] e^{-j\frac{2\pi}{N_2} n_2 k_2} \right] e^{-j\frac{2\pi}{N_1} n_1 k_1}$$

$$\mathbf{V} = \mathbf{W}_{N_2} \mathbf{x}$$

$$\mathbf{V} \in \mathbb{R}^{N_2 \times N_1}$$





$$X[k_1, k_2] = \sum_{n_1=0}^{N_1-1} \left[ \sum_{n_2=0}^{N_2-1} x[n_1, n_2] e^{-j\frac{2\pi}{N_2} n_2 k_2} \right] e^{-j\frac{2\pi}{N_1} n_1 k_1}$$

$$\mathbf{V} = \mathbf{W}_{N_2} \mathbf{x}$$

$$\mathbf{V} \in \mathbb{R}^{N_2 \times N_1}$$

$$\mathbf{X} = \mathbf{V} \mathbf{W}_{N_1}$$

$$\mathbf{X} \in \mathbb{R}^{N_2 \times N_1}$$

$$X[k_1, k_2] = \sum_{n_1=0}^{N_1-1} \left[ \sum_{n_2=0}^{N_2-1} x[n_1, n_2] e^{-j\frac{2\pi}{N_2} n_2 k_2} \right] e^{-j\frac{2\pi}{N_1} n_1 k_1}$$

$$\mathbf{V} = \mathbf{W}_{N_2} \mathbf{x}$$

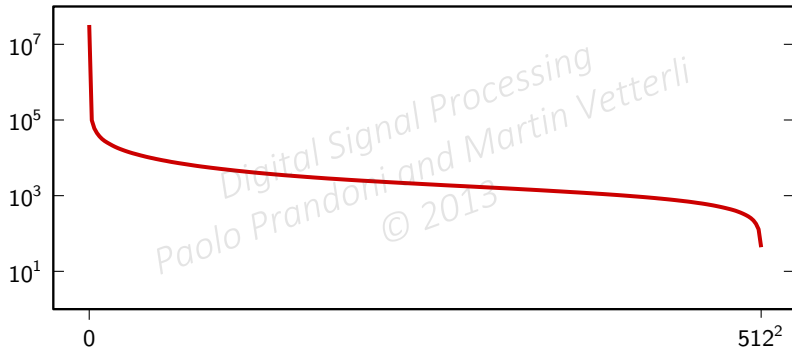
$$\mathbf{V} \in \mathbb{R}^{N_2 \times N_1}$$

$$\mathbf{X} = \mathbf{V} \mathbf{W}_{N_1}$$

$$\mathbf{X} \in \mathbb{R}^{N_2 \times N_1}$$

$$\mathbf{X} = \mathbf{W}_{N_2} \mathbf{x} \mathbf{W}_{N_1}$$

- ▶ try to show the magnitude as an image
- ▶ problem: the range is too big for the grayscale range of paper or screen
- ▶ try to normalize:  $|X'[n_1, n_2]| = |X[n_1, n_2]| / \max |X[n_1, n_2]|$
- ▶ but it doesn't work...



if the image is high dynamic range we need to compress the levels

- ▶ remove flagrant outliers (e.g.  $X[0,0] = \sum \sum x[n_1, n_2]$ )
- ▶ use a nonlinear mapping: e.g.  $y = x^{1/B}$  after normalization ( $x \leq 1$ )

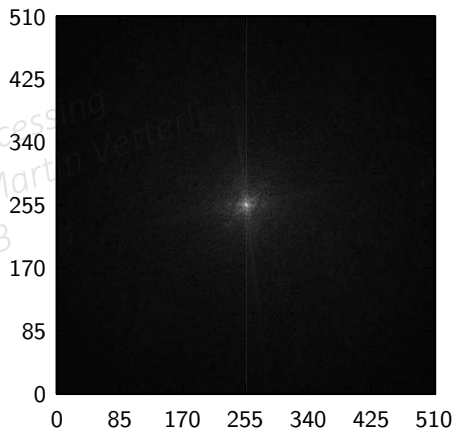
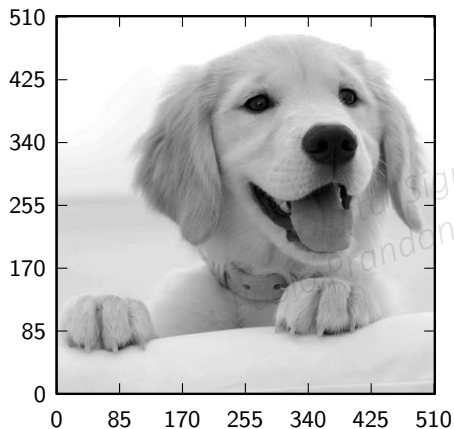
Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

if the image is high dynamic range we need to compress the levels

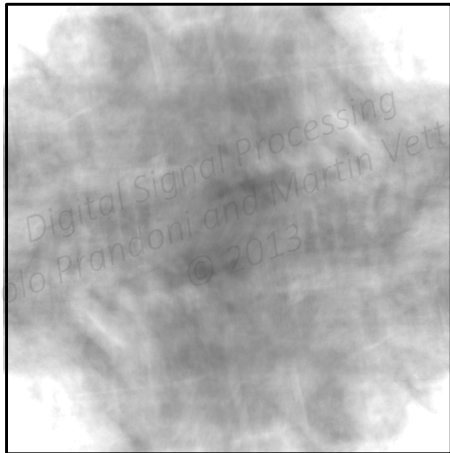
- ▶ remove flagrant outliers (e.g.  $X[0,0] = \sum \sum x[n_1, n_2]$ )
- ▶ use a nonlinear mapping: e.g.  $y = x^{1/3}$  after normalization ( $x \leq 1$ )

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2012

## How does a 2D-DFT look like?

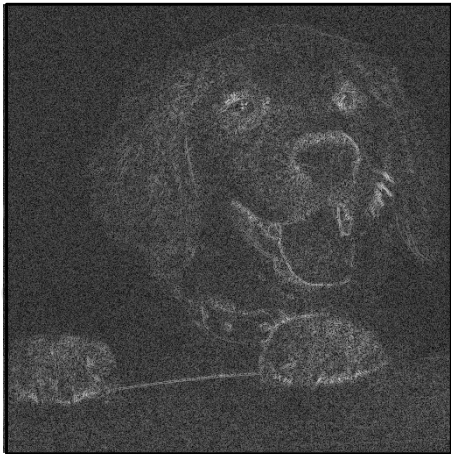


## DFT magnitude doesn't carry much information





DFT phase, on the other hand...



- ▶ most of the information is contained in image's *edges*
- ▶ edges are points of abrupt change in signal's values
- ▶ edges are a “space-domain” feature → not captured by DFT's magnitude
- ▶ phase alignment is important for reproducing edges

END OF MODULE 8.3

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

# Digital Signal Processing

Module 8.4: Filtering

- ▶ Filters for image processing

- ▶ Classification

- ▶ Examples

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ Filters for image processing
- ▶ Classification
- ▶ Examples

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ Filters for image processing
- ▶ Classification
- ▶ Examples

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ linearity
- ▶ *space* invariance
- ▶ impulse response
- ▶ frequency response
- ▶ stability
- ▶ 2D CCDE

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013



- ▶ interesting images contain lots of *semantics*: different information in different areas
- ▶ space-invariant filters process everything in the same way
- ▶ but we should process things differently
  - edges
  - gradients
  - textures
  - ...

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ interesting images contain lots of *semantics*: different information in different areas
- ▶ space-invariant filters process everything in the same way
- ▶ but we should process things differently
  - edges
  - gradients
  - textures
  - ...

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ interesting images contain lots of *semantics*: different information in different areas
- ▶ space-invariant filters process everything in the same way
- ▶ but we should process things differently
  - edges
  - gradients
  - textures
  - ...

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ interesting images contain lots of *semantics*: different information in different areas
- ▶ space-invariant filters process everything in the same way
- ▶ but we should process things differently
  - edges
  - gradients
  - textures
  - ...

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ interesting images contain lots of *semantics*: different information in different areas
- ▶ space-invariant filters process everything in the same way
- ▶ but we should process things differently
  - edges
  - gradients
  - textures
  - ...

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ interesting images contain lots of *semantics*: different information in different areas
- ▶ space-invariant filters process everything in the same way
- ▶ but we should process things differently
  - edges
  - gradients
  - textures
  - ...

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ interesting images contain lots of *semantics*: different information in different areas
- ▶ space-invariant filters process everything in the same way
- ▶ but we should process things differently
  - edges
  - gradients
  - textures
  - ...

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ IIR, FIR

- ▶ causal or noncausal

- ▶ highpass, lowpass, ...

- lowpass → image smoothing

- highpass → enhancement, edge detection

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013



- ▶ IIR, FIR
- ▶ causal or noncausal
- ▶ highpass, lowpass, ...

- lowpass → image smoothing

- highpass → enhancement, edge detection

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ IIR, FIR
- ▶ causal or noncausal
- ▶ highpass, lowpass, ...
  - lowpass → image smoothing
  - highpass → enhancement, edge detection

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ IIR, FIR
- ▶ causal or noncausal
- ▶ highpass, lowpass, ...
  - lowpass → image smoothing
  - highpass → enhancement, edge detection

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ IIR, FIR
- ▶ causal or noncausal
- ▶ highpass, lowpass, ...
  - lowpass → image smoothing
  - highpass → enhancement, edge detection

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ nonlinear phase (edges!)
- ▶ border effects
- ▶ stability: the fundamental theorem of algebra doesn't hold in multiple dimensions!
- ▶ computability

Digital Signal Processing  
paolo Prandoni and Martin Vetterli  
© 2019

- ▶ nonlinear phase (edges!)
- ▶ border effects
- ▶ stability: the fundamental theorem of algebra doesn't hold in multiple dimensions!
- ▶ computability

Digital Signal Processing  
paolo Prandoni and Martin Vetterli  
© 2019

- ▶ nonlinear phase (edges!)
- ▶ border effects
- ▶ stability: the fundamental theorem of algebra doesn't hold in multiple dimensions!
- ▶ computability

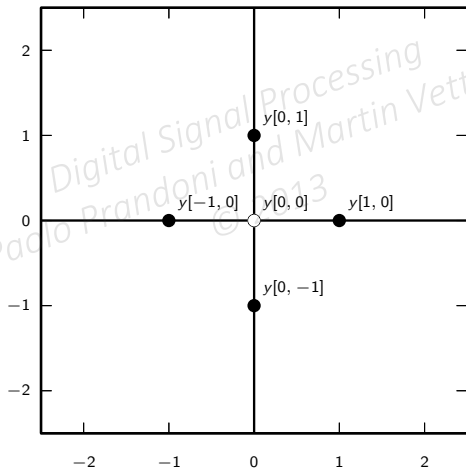
Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ nonlinear phase (edges!)
- ▶ border effects
- ▶ stability: the fundamental theorem of algebra doesn't hold in multiple dimensions!
- ▶ computability

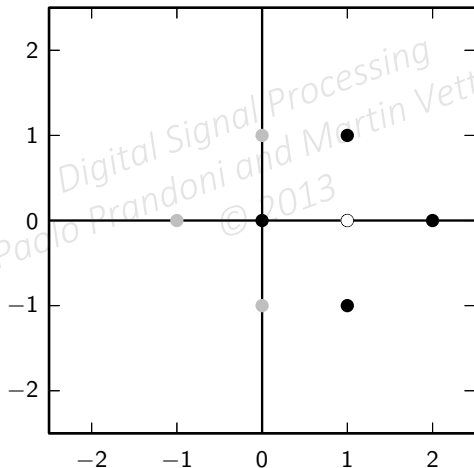
Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013



$$y[n_1, n_2] = a_0 y[n_1 + 1, n_2] + a_1 y[n_1, n_2 - 1] + a_2 y[n_1 - 1, n_2] + a_3 y[n_1, n_2 + 1] + x[n_1, n_2];$$



$$y[n_1, n_2] = a_0 y[n_1 + 1, n_2] + a_1 y[n_1, n_2 - 1] + a_2 y[n_1 - 1, n_2] + a_3 y[n_1, n_2 + 1] + x[n_1, n_2];$$



- ▶ generally zero centered (causality not an issue)  $\Rightarrow$  odd number of taps in both directions

- ▶ per-sample complexity:

- $M_1 M_2$  for nonseparable impulse responses
- $M_1 + M_2$  for separable impulse responses

- ▶ obviously always stable

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ generally zero centered (causality not an issue)  $\Rightarrow$   
odd number of taps in both directions

- ▶ per-sample complexity:

- $M_1 M_2$  for nonseparable impulse responses
- $M_1 + M_2$  for separable impulse responses

- ▶ obviously always stable

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ generally zero centered (causality not an issue)  $\Rightarrow$  odd number of taps in both directions

- ▶ per-sample complexity:

- $M_1 M_2$  for nonseparable impulse responses

- $M_1 + M_2$  for separable impulse responses

- ▶ obviously always stable

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ generally zero centered (causality not an issue)  $\Rightarrow$  odd number of taps in both directions

- ▶ per-sample complexity:

- $M_1 M_2$  for nonseparable impulse responses
- $M_1 + M_2$  for separable impulse responses

- ▶ obviously always stable

- ▶ generally zero centered (causality not an issue)  $\Rightarrow$  odd number of taps in both directions
- ▶ per-sample complexity:
  - $M_1 M_2$  for nonseparable impulse responses
  - $M_1 + M_2$  for separable impulse responses
- ▶ obviously always stable

$$y[n_1, n_2] = \frac{1}{(2N+1)^2} \sum_{k_1=-N}^N \sum_{k_2=-N}^N x[n_1 - k_1, n_2 - k_2]$$

$$h[n_1, n_2] = \frac{1}{(2N+1)^2} \text{rect}\left(\frac{n_1}{2N}, \frac{n_2}{2N}\right)$$



$$y[n_1, n_2] = \frac{1}{(2N+1)^2} \sum_{k_1=-N}^N \sum_{k_2=-N}^N x[n_1 - k_1, n_2 - k_2]$$

$$h[n_1, n_2] = \frac{1}{(2N+1)^2} \text{rect}\left(\frac{n_1}{2N}, \frac{n_2}{2N}\right)$$

$$h[n_1, n_2] = \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013



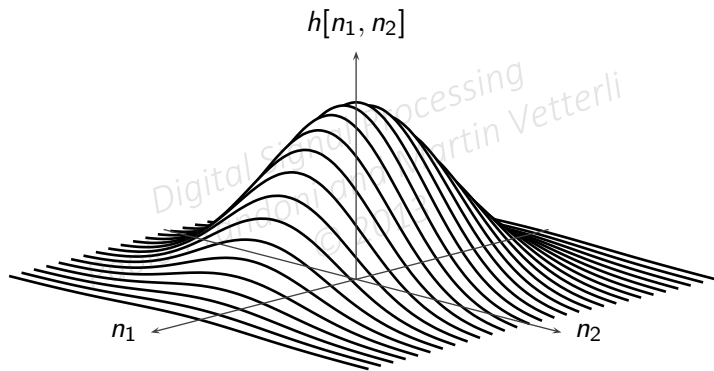
$11 \times 11$  MA

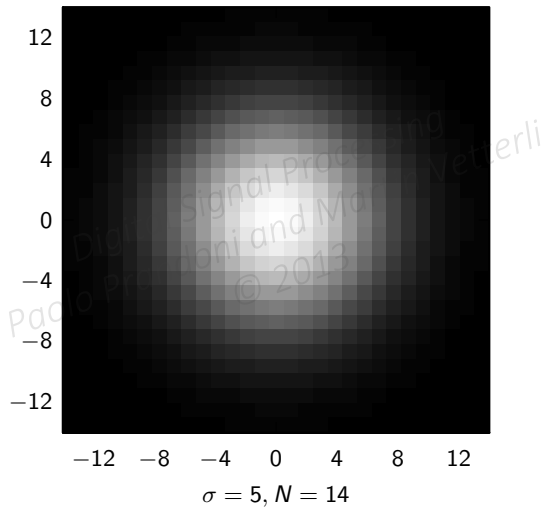


$51 \times 51$  MA

$$h[n_1, n_2] = \frac{1}{2\pi\sigma^2} e^{-\frac{n_1^2 + n_2^2}{2\sigma^2}}, \quad |n_1, n_2| < N$$

with  $N \approx 3\sigma$







$\sigma = 1.8, 11 \times 11$  blur



$\sigma = 8.7, 51 \times 51$  blur

approximation of the first derivative in the horizontal direction:

$$s_o[n_1, n_2] = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

separability and structure

$$s_o[n_1, n_2] = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \begin{bmatrix} -1 & 0 & 1 \end{bmatrix}$$



approximation of the first derivative in the horizontal direction:

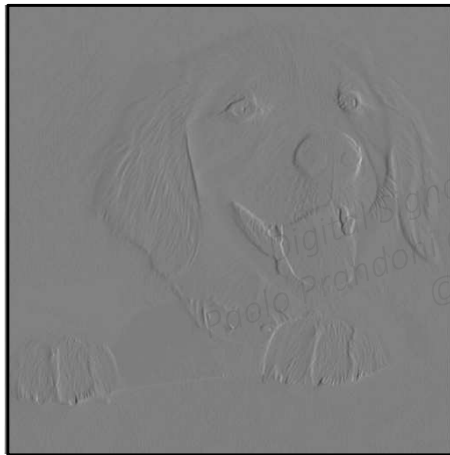
$$s_o[n_1, n_2] = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

separability and structure:

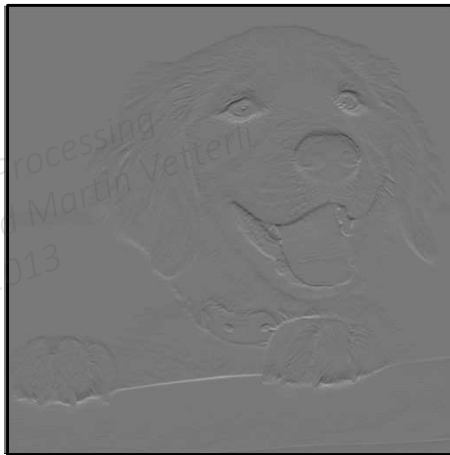
$$s_o[n_1, n_2] = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \begin{bmatrix} -1 & 0 & 1 \end{bmatrix}$$

approximation of the first derivative in the vertical direction:

$$s_v[n_1, n_2] = \begin{bmatrix} -1 & -2 & 1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$



horizontal Sobel filter



vertical Sobel filter

approximation for the square magnitude of the gradient:

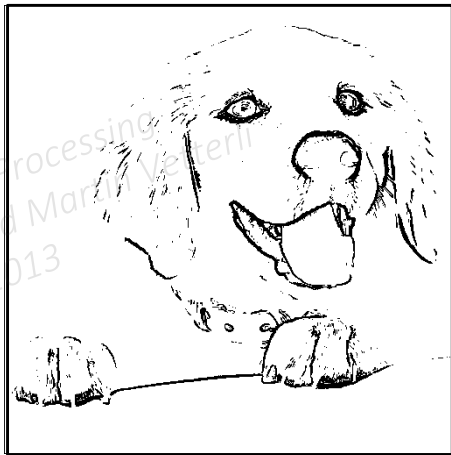
$$|\nabla x[n_1, n_2]|^2 = |s_o[n_1, n_2] * x[n_1, n_2]|^2 + |s_v[n_1, n_2] * x[n_1, n_2]|^2$$

(“operator” because it’s nonlinear)

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013



Sobel operator



thresholded Sobel operator

Laplacian of a function in continuous-space:

$$\Delta f(t_1, t_2) = \frac{\partial^2 f}{\partial t_1^2} + \frac{\partial^2 f}{\partial t_2^2}$$

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

approximating the Laplacian; start with a Taylor expansion

$$f(t + \tau) = \sum_{n=0}^{\infty} \frac{f^{(n)}(t)}{n!} \tau^n$$

and compute the expansion in  $(t + \tau)$  and  $(t - \tau)$ :

$$f(t + \tau) = f(t) + f'(t)\tau + \frac{1}{2}f''(t)\tau^2$$

$$f(t - \tau) = f(t) - f'(t)\tau + \frac{1}{2}f''(t)\tau^2$$

by rearranging terms:

$$f''(t) = \frac{1}{\tau^2}(f(t - \tau) - 2f(t) + f(t + \tau))$$

which, on the discrete grid, is the FIR  $h[n] = [1 \quad -2 \quad 1]$



summing the horizontal and vertical components:

$$h[n_1, n_2] = \begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$$

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

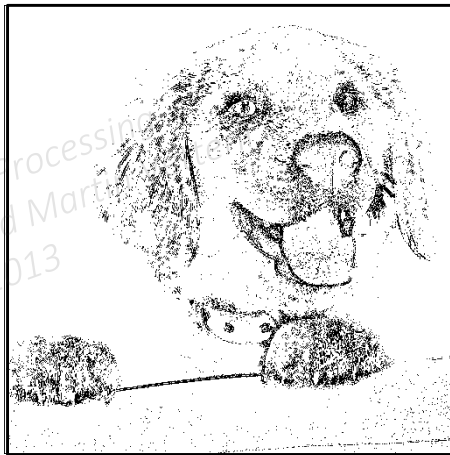
If we use the diagonals too:

$$h[n_1, n_2] = \begin{bmatrix} 1 & 1 & 1 \\ 1 & -8 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013



Laplacian operator



thresholded Laplacian operator

END OF MODULE 8.4

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

# Digital Signal Processing

Module 8.5: Image Compression

- ▶ Redundancy in natural images

- ▶ Image coding ingredients

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ Redundancy in natural images
- ▶ Image coding ingredients

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ consider all possible  $256 \times 256$ , 8bpp images

- ▶ each image is 524,288 bits

- ▶ total number of possible images:  $2^{524,288} \approx 10^{157,826}$

- ▶ number of atoms in the universe:  $10^{82}$

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013



- ▶ consider all possible  $256 \times 256$ , 8bpp images
- ▶ each image is 524,288 bits

- ▶ total number of possible images:  $2^{524,288} \approx 10^{157,826}$
- ▶ number of atoms in the universe:  $10^{82}$

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2004

- ▶ consider all possible  $256 \times 256$ , 8bpp images
- ▶ each image is 524,288 bits
- ▶ total number of possible images:  $2^{524,288} \approx 10^{157,826}$
- ▶ number of atoms in the universe:  $10^{82}$

- ▶ consider all possible  $256 \times 256$ , 8bpp images
- ▶ each image is 524,288 bits
- ▶ total number of possible images:  $2^{524,288} \approx 10^{157,826}$
- ▶ number of atoms in the universe:  $10^{82}$

## Another thought experiment

- ▶ take *all* images in the world and list them in an “encyclopedia of images”
- ▶ to indicate an image, simply give its number
- ▶ on the Internet:  $M = 50$  billion
- ▶ raw encoding: 524,288 bits per image
- ▶ enumeration-based encoding:  $\log_2 M \approx 36$  bits per image
- ▶ (of course, side information is HUGE)

## Another thought experiment

- ▶ take *all* images in the world and list them in an “encyclopedia of images”
- ▶ to indicate an image, simply give its number
- ▶ on the Internet:  $M = 50$  billion
- ▶ raw encoding: 524,288 bits per image
- ▶ enumeration-based encoding:  $\log_2 M \approx 36$  bits per image
- ▶ (of course, side information is HUGE)

## Another thought experiment

- ▶ take *all* images in the world and list them in an “encyclopedia of images”
- ▶ to indicate an image, simply give its number
- ▶ on the Internet:  $M = 50$  billion
- ▶ raw encoding: 524,288 bits per image
- ▶ enumeration-based encoding:  $\log_2 M \approx 36$  bits per image
- ▶ (of course, side information is HUGE)

## Another thought experiment

- ▶ take *all* images in the world and list them in an “encyclopedia of images”
- ▶ to indicate an image, simply give its number
- ▶ on the Internet:  $M = 50$  billion
- ▶ raw encoding: 524,288 bits per image
- ▶ enumeration-based encoding:  $\log_2 M \approx 36$  bits per image
- ▶ (of course, side information is HUGE)

## Another thought experiment

- ▶ take *all* images in the world and list them in an “encyclopedia of images”
- ▶ to indicate an image, simply give its number
- ▶ on the Internet:  $M = 50$  billion
- ▶ raw encoding: 524,288 bits per image
- ▶ enumeration-based encoding:  $\log_2 M \approx 36$  bits per image
- ▶ (of course, side information is HUGE)



## Another thought experiment

- ▶ take *all* images in the world and list them in an “encyclopedia of images”
- ▶ to indicate an image, simply give its number
- ▶ on the Internet:  $M = 50$  billion
- ▶ raw encoding: 524,288 bits per image
- ▶ enumeration-based encoding:  $\log_2 M \approx 36$  bits per image
- ▶ (of course, side information is HUGE)

Another approach:

- ▶ exploit “physical” redundancy
- ▶ allocate bits for things that matter (e.g. edges)
- ▶ use psychovisual experiments to find out what matters
- ▶ some information is discarded: *lossy* compression

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

Another approach:

- ▶ exploit “physical” redundancy
- ▶ allocate bits for things that matter (e.g. edges)
- ▶ use psychovisual experiments to find out what matters
- ▶ some information is discarded: *lossy* compression

Another approach:

- ▶ exploit “physical” redundancy
- ▶ allocate bits for things that matter (e.g. edges)
- ▶ use psychovisual experiments to find out what matters
- ▶ some information is discarded: *lossy* compression

Another approach:

- ▶ exploit “physical” redundancy
- ▶ allocate bits for things that matter (e.g. edges)
- ▶ use psychovisual experiments to find out what matters
- ▶ some information is discarded: *lossy* compression

- ▶ compressing at block level

- ▶ using a suitable transform (i.e., a change of basis)

- ▶ smart quantization

- ▶ entropy coding

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ compressing at block level
- ▶ using a suitable transform (i.e., a change of basis)
- ▶ smart quantization
- ▶ entropy coding

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ compressing at block level
- ▶ using a suitable transform (i.e., a change of basis)
- ▶ smart quantization
- ▶ entropy coding

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013



- ▶ compressing at block level
- ▶ using a suitable transform (i.e., a change of basis)
- ▶ smart quantization
- ▶ entropy coding

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ **reduce number bits per pixel**

- ▶ equivalent to coarser quantization

- ▶ in the limit, 1bpp

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ reduce number bits per pixel
- ▶ equivalent to coarser quantization
- ▶ in the limit, 1bpp

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ reduce number bits per pixel
- ▶ equivalent to coarser quantization
- ▶ in the limit, 1bpp



- ▶ divide the image in blocks

- ▶ code the average value with 8 bits

- ▶  $3 \times 3$  blocks at 8 bits per block gives less than 1bpp

Digital Signal Processing  
paolo Prandoni and Martin Vetterli  
© 2013

- ▶ divide the image in blocks
- ▶ code the average value with 8 bits
- ▶  $3 \times 3$  blocks at 8 bits per block gives less than 1bpp

Digital Signal Processing  
paolo Prandoni and Martin Vetterli  
© 2013

- ▶ divide the image in blocks
- ▶ code the average value with 8 bits
- ▶  $3 \times 3$  blocks at 8 bits per block gives less than 1bpp



- ▶ **exploit the local spatial correlation**

- ▶ compress remote regions independently

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

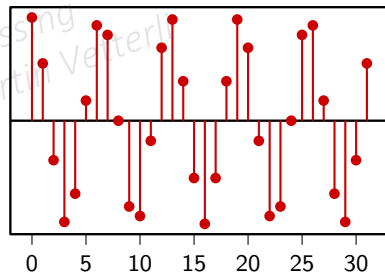


- ▶ exploit the local spatial correlation
- ▶ compress remote regions independently

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

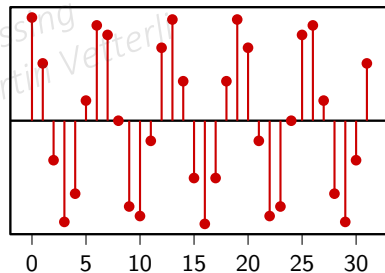
A simple example:

- ▶ take a DT signal, assume  $R$  bits per sample
- ▶ storing the signal requires  $NR$  bits
- ▶ now you take the DFT and it looks like this
- ▶ in theory, we can just code the two nonzero DFT coefficients!



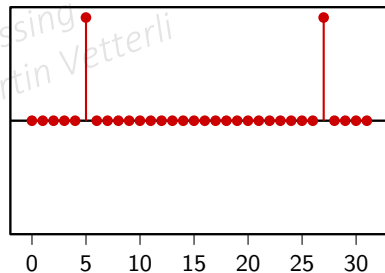
A simple example:

- ▶ take a DT signal, assume  $R$  bits per sample
- ▶ storing the signal requires  $NR$  bits
- ▶ now you take the DFT and it looks like this
- ▶ in theory, we can just code the two nonzero DFT coefficients!



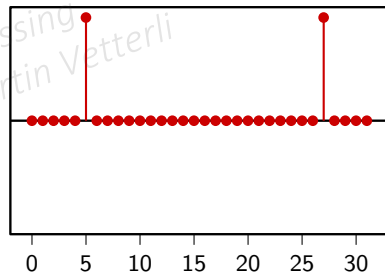
A simple example:

- ▶ take a DT signal, assume  $R$  bits per sample
- ▶ storing the signal requires  $NR$  bits
- ▶ now you take the DFT and it looks like this
- ▶ in theory, we can just code the two nonzero DFT coefficients!



A simple example:

- ▶ take a DT signal, assume  $R$  bits per sample
- ▶ storing the signal requires  $NR$  bits
- ▶ now you take the DFT and it looks like this
- ▶ in theory, we can just code the two nonzero DFT coefficients!



Ideally, we would like a transform that:

- ▶ captures the important features of an image block in a few coefficients
- ▶ is efficient to compute
- ▶ answer: the Discrete Cosine Transform

Digital Signal Processing  
Polo Prandoni and Martin Vetterli  
© 2013

Ideally, we would like a transform that:

- ▶ captures the important features of an image block in a few coefficients
- ▶ is efficient to compute
- ▶ answer: the Discrete Cosine Transform

Ideally, we would like a transform that:

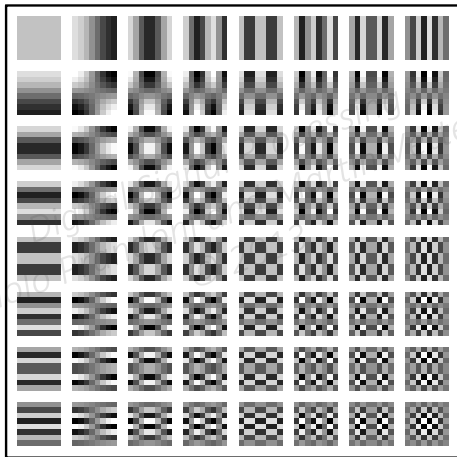
- ▶ captures the important features of an image block in a few coefficients
- ▶ is efficient to compute
- ▶ answer: the Discrete Cosine Transform



$$C[k_1, k_2] = \sum_{n_1=0}^{N-1} \sum_{n_2=0}^{N-1} x[n_1, n_2] \cos \left[ \frac{\pi}{N} \left( n_1 + \frac{1}{2} \right) k_1 \right] \cos \left[ \frac{\pi}{N} \left( n_2 + \frac{1}{2} \right) k_2 \right]$$

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

## DCT basis vectors for an $8 \times 8$ image



- ▶ **deadzone**

- ▶ variable step (fine to coarse)

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

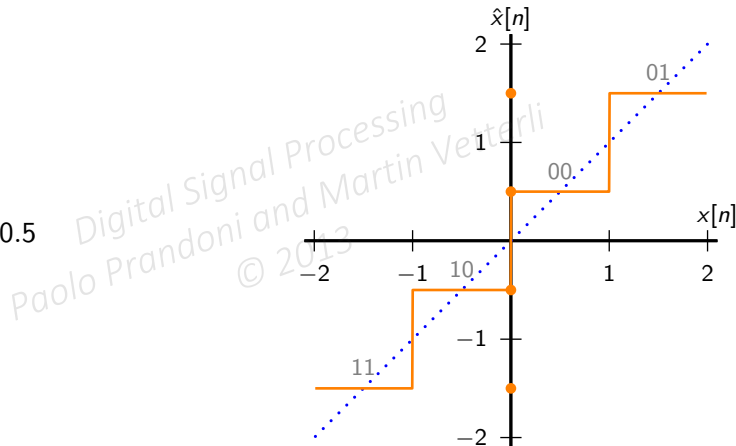
- ▶ deadzone

- ▶ variable step (fine to coarse)

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

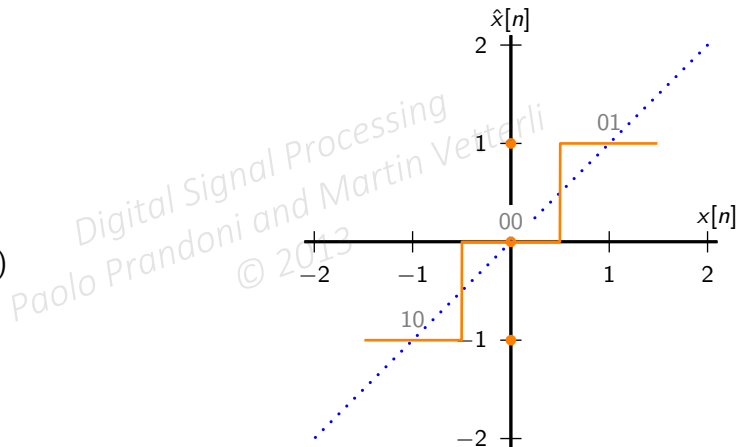
Standard quantization:

$$\hat{x} = \text{floor}(x) + 0.5$$



Deadzone quantization:

$$\hat{x} = \text{round}(x)$$



- ▶ minimize the effort to encode a certain amount of information
- ▶ associate short symbols to frequent values and vice-versa
- ▶ if it sounds familiar it's because it is...

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ minimize the effort to encode a certain amount of information
- ▶ associate short symbols to frequent values and vice-versa
- ▶ if it sounds familiar it's because it is...

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013



- ▶ minimize the effort to encode a certain amount of information
- ▶ associate short symbols to frequent values and vice-versa
- ▶ if it sounds familiar it's because it is...

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

A	•	■			
B	■	•	•	•	
C	■	•	■	•	
D	■	•	•		
E	•				
F	•	•	■	•	
G	■	■	■	•	
H	•	•	•	•	
I	•	•			
J	•	■	■	■	■
K	■	•	■	■	
L	•	■	•	•	
M	■	■	■		
N	■	•			
O	■	■	■	■	
P	•	■	■	■	•
Q	■	■	■	•	■
R	•	■	■	•	
S	•	•	•		
T	■				

U	•	•	■		
V	•	•	•	■	
W	•	■	■	■	
X	■	•	•	■	
Y	■	•	■	■	■
Z	■	■	■	•	•

1	•	■	■	■	■	■
2	•	■	■	■	■	■
3	•	•	•	■	■	■
4	•	•	•	•	■	
5	•	•	•	•	•	
6	■	•	•	•	•	
7	■	■	■	•	•	•
8	■	■	■	■	•	•
9	■	■	■	■	■	•
0	■	■	■	■	■	■

Digital Signal Processing  
 Paolo Prandoni and Martin Vetterli  
 © 2013

END OF MODULE 8.5

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

# Digital Signal Processing

## Module 8.6: The JPEG Compression Algorithm

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ compressing at block level
- ▶ using a suitable transform (i.e., a change of basis)
- ▶ smart quantization
- ▶ entropy coding

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ split image into  $8 \times 8$  non-overlapping blocks
- ▶ using a suitable transform (i.e., a change of basis)
- ▶ smart quantization
- ▶ entropy coding

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ split image into  $8 \times 8$  non-overlapping blocks
- ▶ compute the DCT of each block
- ▶ smart quantization
- ▶ entropy coding

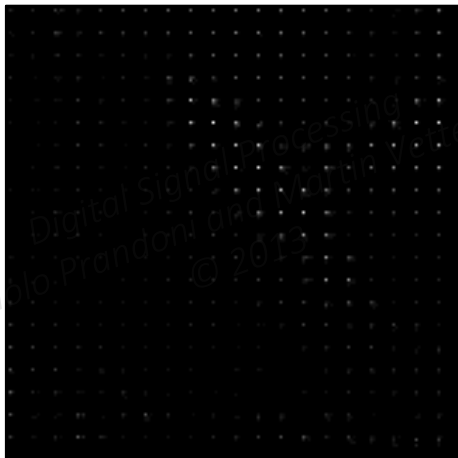
Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ split image into  $8 \times 8$  non-overlapping blocks
- ▶ compute the DCT of each block
- ▶ quantize DCT coefficients according to psychovisually-tuned tables
- ▶ entropy coding

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

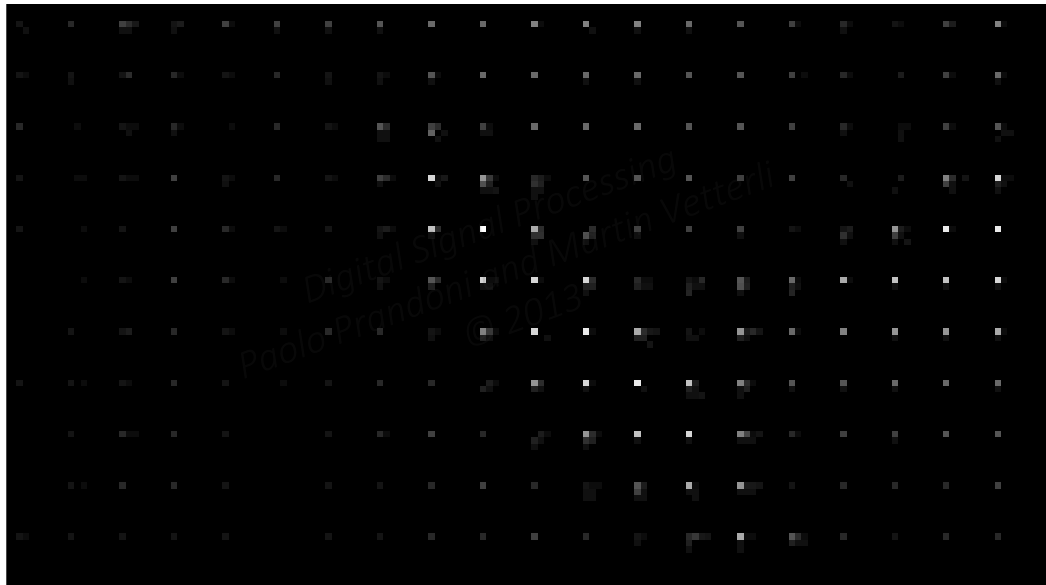


- ▶ split image into  $8 \times 8$  non-overlapping blocks
- ▶ compute the DCT of each block
- ▶ quantize DCT coefficients according to psychovisually-tuned tables
- ▶ run-length encoding and Huffman coding



Digital Signal Processing  
Pablo Prandoni and Martin Vetterli  
© 2013

## DCT coefficients of image blocks (detail)



- ▶ most coefficients are negligible → captured by the deadzone
- ▶ some coefficients are more important than others
- ▶ find out the critical coefficients by experimentation
- ▶ allocate more bits (or, equivalently, finer quantization levels) to the most important coefficients

- ▶ most coefficients are negligible  $\rightarrow$  captured by the deadzone
- ▶ some coefficients are more important than others
- ▶ find out the critical coefficients by experimentation
- ▶ allocate more bits (or, equivalently, finer quantization levels) to the most important coefficients

- ▶ most coefficients are negligible  $\rightarrow$  captured by the deadzone
- ▶ some coefficients are more important than others
- ▶ find out the critical coefficients by experimentation
- ▶ allocate more bits (or, equivalently, finer quantization levels) to the most important coefficients

- ▶ most coefficients are negligible  $\rightarrow$  captured by the deadzone
- ▶ some coefficients are more important than others
- ▶ find out the critical coefficients by experimentation
- ▶ allocate more bits (or, equivalently, finer quantization levels) to the most important coefficients

$$\hat{c}[k_1, k_2] = \text{round}(c[k_1, k_2]/Q[k_1, k_2])$$

$$Q = \begin{bmatrix} 16 & 11 & 10 & 16 & 24 & 40 & 51 & 61 \\ 12 & 12 & 14 & 19 & 26 & 58 & 60 & 55 \\ 14 & 13 & 16 & 24 & 40 & 57 & 69 & 56 \\ 14 & 17 & 22 & 29 & 51 & 87 & 80 & 62 \\ 18 & 22 & 37 & 56 & 68 & 109 & 103 & 77 \\ 24 & 35 & 55 & 64 & 81 & 104 & 113 & 92 \\ 49 & 64 & 78 & 87 & 103 & 121 & 120 & 101 \\ 72 & 92 & 95 & 98 & 112 & 100 & 103 & 99 \end{bmatrix}$$





uniform



tuned



uniform



tuned

Visual Signal Processing  
Paolo F. Feltoni and Martin Vetterli  
© 2013

- ▶ most coefficients are small, decreasing with index

- ▶ use zigzag scan to maximize ordering

- ▶ quantization will create long series of zeros

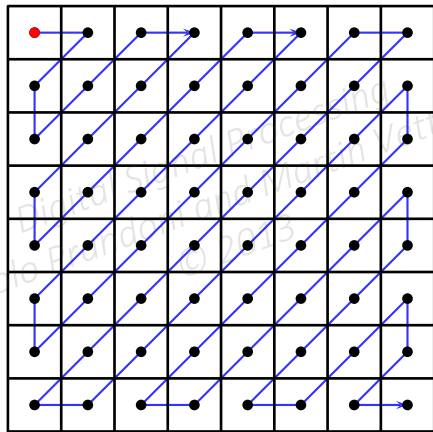
Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ most coefficients are small, decreasing with index
- ▶ use zigzag scan to maximize ordering
- ▶ quantization will create long series of zeros

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

- ▶ most coefficients are small, decreasing with index
- ▶ use zigzag scan to maximize ordering
- ▶ quantization will create long series of zeros

Digital Signal Processing  
Paolo Prasadoni and Martin Vetterli  
© 2013



$$\begin{bmatrix} 100 & -60 & 0 & 6 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 13 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

$$\begin{bmatrix} 100 & -60 & 0 & 6 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 13 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 13 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$



Each nonzero value is encoded as the triple

$$[(r, s), c]$$

- ▶  $r$  is the *runlength* i.e. the number of zeros before the current value
- ▶  $s$  is the *size* i.e. the number of bits needed to encode the value
- ▶  $c$  is the actual value
- ▶  $(0, 0)$  indicates that from now on it's only zeros (end of block)

Each nonzero value is encoded as the triple

$$[(r, s), c]$$

- ▶  $r$  is the *runlength* i.e. the number of zeros before the current value
- ▶  $s$  is the *size* i.e. the number of bits needed to encode the value
- ▶  $c$  is the actual value
- ▶  $(0, 0)$  indicates that from now on it's only zeros (end of block)

Each nonzero value is encoded as the triple

$$[(r, s), c]$$

- ▶  $r$  is the *runlength* i.e. the number of zeros before the current value
- ▶  $s$  is the *size* i.e. the number of bits needed to encode the value
- ▶  $c$  is the actual value
- ▶  $(0, 0)$  indicates that from now on it's only zeros (end of block)

Each nonzero value is encoded as the triple

$$[(r, s), c]$$

- ▶  $r$  is the *runlength* i.e. the number of zeros before the current value
- ▶  $s$  is the *size* i.e. the number of bits needed to encode the value
- ▶  $c$  is the actual value
- ▶  $(0, 0)$  indicates that from now on it's only zeros (end of block)

$$\begin{bmatrix} 100 & -60 & 0 & 6 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 13 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

$[(0, 7), 100], [(0, 6), -60], [(4, 3), 6], [(3, 4), 13], [(8, 1), -1], [(0, 0)]$

$$\begin{bmatrix} 100 & -60 & 0 & 6 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 13 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

$[(0, 7), 100], [(0, 6), -60], [(4, 3), 6], [(3, 4), 13], [(8, 1), -1], [(0, 0)]$

- ▶ by design,  $(r, s) \in \mathcal{A}$  with  $|\mathcal{A}| = 256$

- ▶ in theory, 8 bits per pair

- ▶ some pairs are much more common than others

- ▶ a lot of space can be saved by being smart

Digital Signal Processing  
paolo Prandoni and Martin Vetterli  
© 2013

- ▶ by design,  $(r, s) \in \mathcal{A}$  with  $|\mathcal{A}| = 256$

- ▶ in theory, 8 bits per pair

- ▶ some pairs are much more common than others

- ▶ a lot of space can be saved by being smart

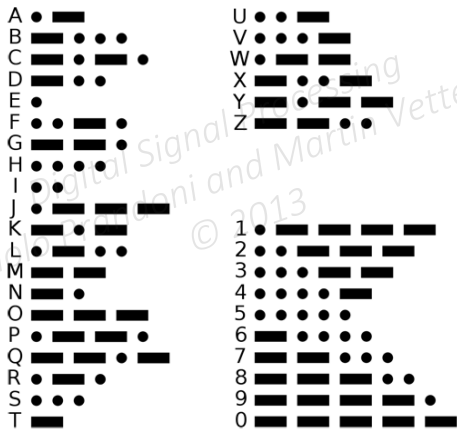
Digital Signal Processing  
paolo Prandoni and Martin Vetterli  
© 2013



- ▶ by design,  $(r, s) \in \mathcal{A}$  with  $|\mathcal{A}| = 256$
- ▶ in theory, 8 bits per pair
- ▶ some pairs are much more common than others!
- ▶ a lot of space can be saved by being smart

- ▶ by design,  $(r, s) \in \mathcal{A}$  with  $|\mathcal{A}| = 256$
- ▶ in theory, 8 bits per pair
- ▶ some pairs are much more common than others!
- ▶ a lot of space can be saved by being smart

great idea: shorter binary sequences for common symbols



A	• ■■	U	• • ■■
B	■■■ • • •	V	• • • ■■
C	■■■ • ■■ •	W	• ■■ ■■
D	■■■ • •	X	■■■ • • ■■
E	•	Y	■■■ • ■■ ■■
F	• • ■■ •	Z	■■■ ■■ • •
G	■■■ ■■ •		
H	• • • •		
I	• • ■■ ■■ ■■		
J	• ■■ ■■ ■■		
K	■■■ • ■■		
L	• ■■ • •		
M	■■■ ■■		
N	■■■ •		
O	■■■ ■■ ■■		
P	• ■■ ■■ •		
Q	■■■ ■■ • ■■		
R	• ■■ •		
S	• • •		
T	■■■		
		1	• ■■ ■■ ■■ ■■
		2	• • ■■ ■■ ■■
		3	• • • ■■ ■■
		4	• • • • ■■
		5	• • • • •
		6	■■■ • • • •
		7	■■■ ■■ • • •
		8	■■■ ■■ ■■ • •
		9	■■■ ■■ ■■ ■■ •
		0	■■■ ■■ ■■ ■■ ■■

however: if symbols have different lengths, we must know how to parse them!

- ▶ in English, spaces separate words → extra symbol (wasteful)
- ▶ in Morse code, pauses separate letters and words (wasteful)
- ▶ can we do away with separators?

however: if symbols have different lengths, we must know how to parse them!

- ▶ in English, spaces separate words → extra symbol (wasteful)
- ▶ in Morse code, pauses separate letters and words (wasteful)
- ▶ can we do away with separators?

however: if symbols have different lengths, we must know how to parse them!

- ▶ in English, spaces separate words → extra symbol (wasteful)
- ▶ in Morse code, pauses separate letters and words (wasteful)
- ▶ can we do away with separators?

- ▶ no valid sequence can be the beginning of another valid sequence
- ▶ can parse a bitstream sequentially with no look-ahead
- ▶ extremely easy to understand graphically...

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

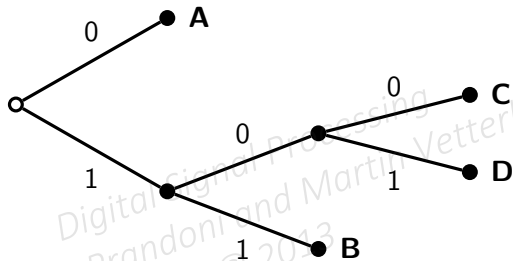
- ▶ no valid sequence can be the beginning of another valid sequence
- ▶ can parse a bitstream sequentially with no look-ahead
- ▶ extremely easy to understand graphically...

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

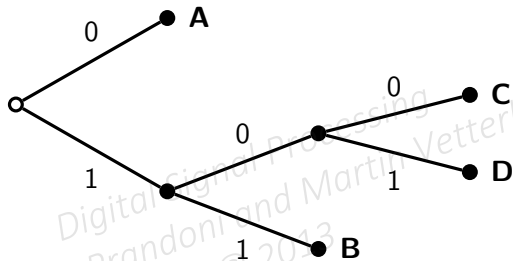


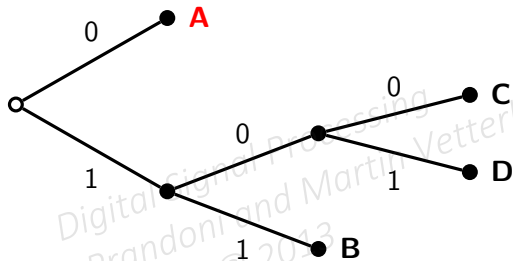
- ▶ no valid sequence can be the beginning of another valid sequence
- ▶ can parse a bitstream sequentially with no look-ahead
- ▶ extremely easy to understand graphically...

Digital Signal Processing  
Paolo P. P. and Martin Vetterli  
© 2013



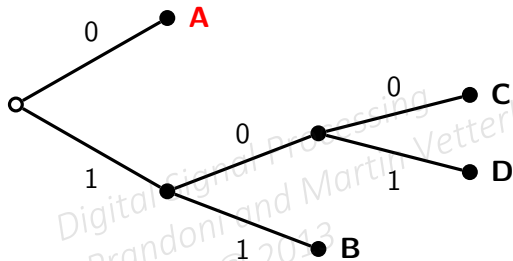
001100110101100





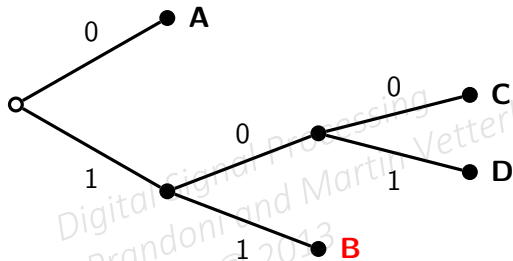
001100110101100

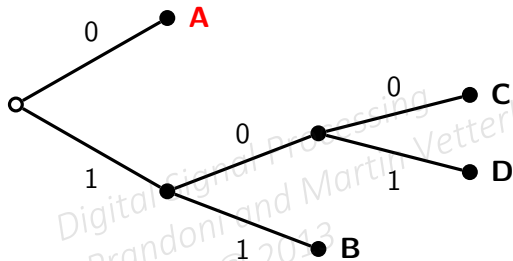
A



001100110101100

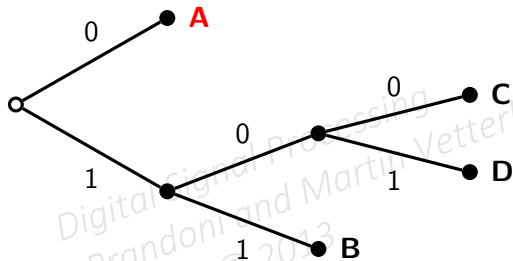
AA





001100110101100

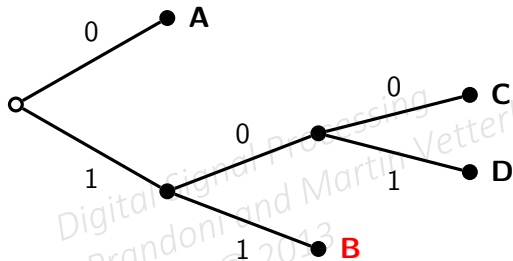
AABA

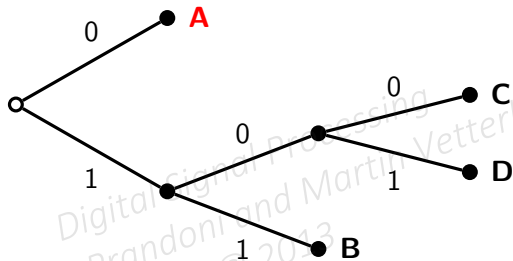


001100110101100

AABAA

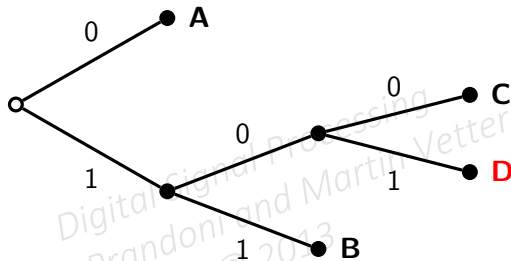


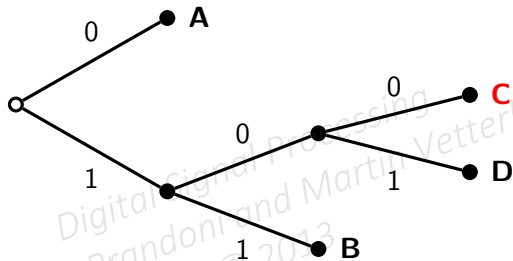




001100110101100

AABAABA





goal: minimize message length

- ▶ assign short sequences to more frequent symbols
- ▶ the Huffman algorithm builds the optimal code for a set of symbol probabilities
- ▶ in JPEG, you can use a “general-purpose” Huffman code or build your own (but then you pay a “side-information” price)

Digital Signal Processing  
© 2013  
Paolo Prandoni and Martin Vetterli

goal: minimize message length

- ▶ assign short sequences to more frequent symbols
- ▶ the Huffman algorithm builds the optimal code for a set of symbol probabilities
- ▶ in JPEG, you can use a “general-purpose” Huffman code or build your own (but then you pay a “side-information” price)

goal: minimize message length

- ▶ assign short sequences to more frequent symbols
- ▶ the Huffman algorithm builds the optimal code for a set of symbol probabilities
- ▶ in JPEG, you can use a “general-purpose” Huffman code or build your own (but then you pay a “side-information” price)

- ▶ four symbols: A, B, C, D

- ▶ probability table:

Digital Signal Processing  
paolo.prandoni and Martin Vetterli  
© 2013

$p(A) = 0.38$	$p(B) = 0.32$
$p(C) = 0.1$	$p(D) = 0.2$



- ▶ four symbols: A, B, C, D
- ▶ probability table:

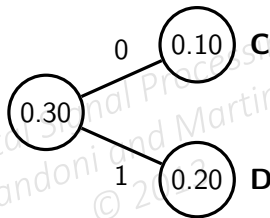
$$p(A) = 0.38$$

$$p(B) = 0.32$$

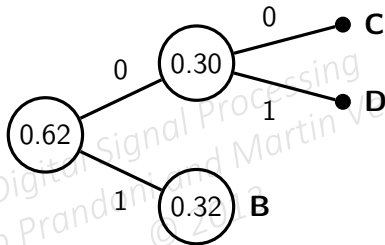
$$p(C) = 0.1$$

$$p(D) = 0.2$$

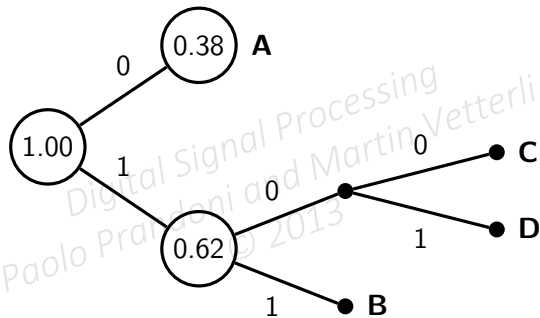
Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013



$$p(A) = 0.38 \quad p(B) = 0.32 \quad p(C) = 0.1 \quad p(D) = 0.2$$



$$p(A) = 0.38 \quad p(B) = 0.32 \quad p(C + D) = 0.3$$



$$p(A) = 0.38 \quad p(B + C + D) = 0.62$$

**END OF MODULE 8.6**

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013

END OF MODULE 8

Digital Signal Processing  
Paolo Prandoni and Martin Vetterli  
© 2013